

When Warm Feelings Harden: Fentanyl Crisis Blame and U.S. Public Opinion on China

Junyang Hu*

April 30, 2026

Abstract

How do Americans form policy preferences toward China when a domestic crisis is externally attributed? Drawing on two nationally representative surveys of U.S. adults fielded in 2024 and 2025, this article introduces the concept of issue-bounded conditionality: the coexistence of coercive and cooperative policy preferences within a single issue domain without extending across areas. Warmth toward China and blame for the fentanyl crisis jointly shape both generalized bilateral posture and issue-level tool preferences, but these two levels of opinion do not move in lockstep. Among respondents who continue to hold relatively warmer views of China, blame is associated with a sanctions-and-cooperation configuration instead of a uniformly confrontational stance. A specificity test spanning eight policy domains shows that this pattern is concentrated in the fentanyl domain, with little evidence of spillover to areas such as cross-Strait tension, economic decoupling, or educational exchange. These findings show that support for issue-specific engagement is most durable when tied to a concrete, accountable objective rather than to abstract appeals to goodwill.

Keywords: U.S.-China relations, public opinion, crisis blame, issue-bounded conditionality

*PAX *sapiens* Foundation. Corresponding author: jhu@paxsapiens.org.

1 Introduction

Over the past decade, claims that China poses a threat to U.S. national interests and to the liberal international order have structured much of Washington’s debate on China, though the precise framing has varied across administrations. One policy domain in which these shifts have had direct and tangible consequences is the U.S. synthetic opioid epidemic. The crisis reaches into nearly every community: nearly one-third of Americans report that they or a family member have experienced opioid addiction, and about one-fifth report knowing someone who died of an overdose (Hu, Baldwin, and Seyle 2025; Saunders, Panchal, and Rudowitz 2026). Fentanyl, a synthetic opioid responsible for over half of U.S. overdose deaths since 2019 and nearly 70% of such deaths in 2023, now exemplifies a transnational public health emergency with geopolitical underpinnings (USAFacts 2025). The People’s Republic of China (PRC) has been identified as a primary source of precursor chemicals used in the illicit production of fentanyl by Mexican cartels.

Despite Beijing’s 2019 regulatory measures to curb fentanyl-related exports, U.S. politicians and pundits have increasingly characterized the PRC’s role as ranging from negligent to deliberately exploitative, with some describing the flow of precursor chemicals as an “asymmetric weapon” (Martínez-Fernández and Harding 2024). In a 2024 report, the House Select Committee on the Chinese Communist Party concluded that the PRC government subsidizes the manufacture and export of illicit fentanyl materials and other synthetic narcotics, fails to prosecute major traffickers, permits open online sales of precursor chemicals, and derives strategic and economic benefit from the crisis (U.S. House Select Committee on the Strategic Competition Between the United States and the Chinese Communist Party 2024). These assessments, in turn, have helped drive congressional proposals to broaden sanctionable categories and designate certain Chinese entities and officials as “foreign opioid traffickers” (Barr 2025).

The executive branch, for its part, has likewise highlighted the growing role of Chinese money-laundering networks in facilitating the operations of Mexico-based transnational criminal organizations (TCOs). Beginning in 2025, the Trump administration folded this line of argument into a broader national-security narrative that portrayed China as an “unusual and extraordinary threat”

and as a source of support and safe haven for fentanyl-linked TCOs (The White House 2025). Together, these developments translated into concrete policy, from the imposition of fentanyl-related tariffs in early 2025 to a bilateral presidential agreement in Busan later that year. In doing so, Washington has embraced a model of transactional *realpolitik* that ties punitive tariffs and technology export controls to stepped-up counternarcotics enforcement.

Although the terms of both Washington’s political debate and its dealings with Beijing have been reshaped, elite-level shifts do not by themselves explain how ordinary Americans translate blame for the opioid crisis into their foreign-policy preferences. Whether ordinary citizens engage in a similar form of transactional reasoning and under what conditions remains an open question. Elite cues on China are often contradictory—simultaneously framing Beijing as a partner in counternarcotics cooperation and as a deliberate enabler of the crisis—and citizens bring their own affective dispositions and issue-specific judgments to bear. Public attitudes toward China remain broadly negative, yet they are far from monolithic. The share describing China as an outright “enemy” fell from 42% to 33% in a single year, and for the first time since 2019, a majority of Americans favored “friendly cooperation and engagement” with Beijing (Huang, Silver, and Clancy 2025; Kafura 2025). The Chicago Council on Global Affairs annual surveys (Kafura 2025) provide the closest national-benchmark comparator for these trends in U.S. attitudes toward China; the 2024–2025 waves analyzed here track the Chicago Council figures on aggregate warmth and cooperation support within 3 percentage points while adding the granularity to decompose tool configurations. Such aggregate figures mask considerable internal differentiation: Americans can hold negative general impressions of a foreign country while supporting cooperation on specific issues, and favorability ratings and policy-domain preferences do not always move in tandem (Holsti 2004; Gries and Crowson 2010; Jin, Dorius, and Xie 2021). More than a thermometer of bilateral tensions, public opinion functions as an independent source of constraint and legitimation in foreign-policy formation.

How this differentiated structure operates when a lethal domestic crisis is attributed to a foreign rival is the puzzle this article addresses. Blame, threat perception, and potential cooperation are

unusually intertwined in the fentanyl case. The resulting opinion dynamics cannot be adequately captured by a single hawk–dove dimension: affective warmth toward China, support for coercive policy tools, and responsiveness to crisis cues need not lie on a single continuum. This premise thus motivates two related questions. First, does the relationship between warmth and foreign-policy posture hold, intensify, or collapse into uniform confrontation when China is blamed for a serious domestic harm? Second, can warmth coexist with endorsement of coercive instruments—and if so, is this coexistence domain-specific or does it generalize across policy areas? Answering these questions requires examining whether crisis blame suppresses, reverses, or reorganizes the relationship between warmth and coercive preferences.

To that end, the article draws on two nationally representative YouGov survey waves in 2024 and 2025. The findings reveal a conditional rather than linear relationship. Warmth toward China is associated with a more dovish general posture, and this association is not eliminated by blaming China for the fentanyl crisis. Among respondents who view China warmly, blame is not associated with a uniform shift toward punishment. Instead, blame is associated with a distinctive mixed configuration—pairing coercive pressure with demands for cooperative engagement—rather than with sanctions alone. An eight-probe specificity test, using seven boundary probes as a negative-control exercise, further confirms that this pattern is domain-specific. The coexistence of warmth and coercive support appears in the fentanyl-cooperation scenario and is largely absent in military, economic, and cultural domains. Warmth toward China coexists with support for coercive instruments when those instruments are tied to a concrete, reciprocal, and issue-specific form of cooperation.

The findings speak to three literatures. The image-theory tradition treats national images as structured by broad relational variables—relative power, perceived threat, cultural status—not issue-specific attributions (Herrmann et al. 1997). Blame intensifying cooperative preferences among those who view China favorably demonstrates that issue-specific cognitions can redirect, not merely reinforce, the policy consequences of affect. The article also draws a distinction between general posture and instrument choice that is undertheorized in research on economic statecraft.

Building on evidence that citizens evaluate sanctions instrumentally, weighing expected consequences rather than expressing generalized hostility (Jentleson 1992; Heinrich, Kobayashi, and Peterson 2017), the results show that preferences for coercive tools can coexist with a cooperative orientation when those tools are perceived as serving a specific bargaining objective. Finally, the article extends research on “issue publics”—subgroups with unusually structured attitudes on salient matters—to the foreign-policy instrument domain (Krosnick 1990; Krosnick and Telhami 1995). Recent work theorizes that such publics form when a pivotal circumstance activates the political relevance of a latent group characteristic (Rossiter and Harden 2026). That warmth-plus-coercion preferences toward China are confined to the fentanyl scenario is consistent with this logic: crisis blame is associated with issue-bounded preference structures that cut across the conventional hawk–dove continuum.

Section 2 reviews the relevant literature and develops the theoretical framework. Section 3 introduces the data, measures, and analytical strategy. Section 4 presents the empirical findings, moving from general posture (Section 4.1) to tool configuration (Section 4.2) and the scenario-based specificity test (Section 4.3). Section 5 discusses implications for the study of U.S.-China relations and democratic foreign-policy formation.

2 Literature and Theory

2.1 From Threat Perception to Issue-Level Reasoning

When domestic crises intersect with foreign actors, threat perception and emotional arousal should dominate public reasoning about foreign policy. Appraisal theories of emotion hold that anger and fear—discrete affective states triggered by perceived threat—generate systematically different risk assessments and policy preferences (Marcus, Neuman, and MacKuen 2000; Lerner and Keltner 2001; Herrmann 2017; Renshon, Lee, and Tingley 2017). In the foreign policy domain, anxiety sometimes prompts risk aversion and impairs cognitive performance, while anger fosters a desire for retributive punishment that can override established partisan leanings (**brader_what_2008**;

Huddy et al. 2005; Wayne 2023). In cases of high-stakes conflict, reminders of external aggression often lead to increased support for hardline policies regardless of whether the specific emotion felt is fear or anger (Kupatadze and Zeitzoff 2021). Similarly, when citizens perceive a foreign actor as both powerful and threatening, adversary-image schemata can also activate, channeling preferences toward confrontation (Herrmann, Tetlock, and Visser 1999). Elite cues reinforce these tendencies. When bipartisan consensus narrows the available frames—as has occurred with “tough on China” rhetoric in U.S. politics—the cue space constricts, and citizens face few elite signals that might encourage more differentiated reasoning (Zaller 1992). Together, these frameworks predict a unidimensional hawkish shift. The more negative the affect toward the rival, the more confrontational the public’s preferences should be across all policy dimensions. Yet this prediction sits uneasily with a recurrent empirical pattern. Some citizens simultaneously endorse punitive measures *and* domain-specific cooperation with the same adversary—a configuration that purely affect-driven models do not capture.

While these frameworks powerfully explain the general hawkish tilt in crisis contexts, they predict a uniformity of response that the empirical record increasingly contradicts. Even within the elite-cueing tradition, recent work shows that cue effects weaken when citizens possess independent policy information or when partisan polarization channels processing through pre-existing attitudes over elite frames (Bullock 2011; Guisinger and Saunders 2017). A more fundamental challenge comes from research demonstrating that public foreign policy attitudes are not volatile reactions to elite signals but hierarchically organized belief systems in which general orientations—rooted in values, dispositions, and social identifications—constrain more specific policy preferences (Hurwitz and Peffley 1987). Citizens form coherent foreign policy attitudes even without elite guidance, suggesting that preference structures are built from the bottom up rather than imposed from the top down (Kertzer and Zeitzoff 2017). The same basic personal values that guide daily choices also structure foreign-policy orientations: conservation values predict militant internationalism while universalism predicts cooperative internationalism, a pattern that holds even among low-knowledge respondents (Rathbun et al. 2016). Across decades and issue domains, aggregate public opinion

exhibits a degree of stability and internal coherence that cannot be reduced to a simple hawk–dove continuum (Page and Shapiro 1992; Kertzer 2023).

This multidimensionality is visible in the U.S.-China case too, where public attitudes are constructed beyond a single favorable–unfavorable metric (Li 2021). Despite bipartisan consensus that China “dodges” international responsibility based on its domestic governance and strategic weight, ideology and partisanship produce sharp cleavages over China policy (Gries and Crowson 2010; Gries 2014; Li and Ye 2015). The pattern sharpens when specific policy domains are examined. Americans can simultaneously view bilateral trade as beneficial and endorse the trade war as legitimate, with the two judgments driven by distinct factors—general trade attitudes shaped by cultural perceptions of China, trade-war support predicted primarily by partisan identification (Jin, Dorius, and Xie 2021). The coexistence has limits, though. Experimental data suggests that perceived economic opportunity improves American attitudes toward democratic partners but fails to moderate hostility toward authoritarian states, where judgments of China’s political culture dominate (Wick 2014). Attitudes are structured and internally differentiated. But structure along what dimensions? Specifically, is the relevant structure organized around the *objectives* of foreign policy action, the *instruments* through which action is applied, or the *issue domains* on which action is taken? The “structured opinion” finding opens this dimensionality question without answering it.

One influential answer identifies the principal policy objective (PPO) as the primary axis of differentiation. Jentleson’s “pretty prudent public” thesis demonstrated that American support for the use of military force varies systematically by the stated objective of intervention. Restraining an external aggressor or humanitarian intervention both receive substantially higher support than internal political change (Jentleson 1992; Jentleson and Britton 1998). This objective-sensitivity proved robust across the post-Cold War period, surviving changes in threat environment and partisan alignment. More recent work extends the insight beyond military force to economic statecraft, showing that citizens’ trade preferences are shaped by geopolitical alignment with trading partners—preferring allies over adversaries due to security concerns—rather than being determined solely by material self-interest (Carnegie and Gaikwad 2022). The public applies conditional judgment

based on the perceived purpose of foreign policy action, not a blanket hawkish or dovish disposition. Yet the PPO framework, for all its explanatory power, implicitly treats alternative instruments as interchangeable means to a shared end. A citizen who endorses “restraining an aggressor” may nonetheless evaluate sanctions, diplomatic pressure, and military action very differently—even when all three serve the same objective. Whether citizens treat instruments as interchangeable conditional on endorsing a given objective remains underspecified.

Recent work on public support for specific coercive tools suggests they do not. Even holding objectives constant, instruments carry independent weight in public evaluation. A growing experimental literature examines the conditions under which mass publics support coercive versus cooperative instruments in bilateral relationships (**tomz_domestic_2007**; Tomz and Weeks 2013; Heinrich, Kobayashi, and Peterson 2017; Frye 2019). Survey-experimental evidence shows that citizens support sanctions more when they expect the measures to deter future misbehavior by the target state; support declines as anticipated domestic economic costs to the sender increase (Heinrich, Kobayashi, and Peterson 2017). Related work shows that human rights framing and target-population salience further condition public support for sanctions, with tailored sanctions targeting leadership tending to attract greater backing than comprehensive ones (McLean and Roblyer 2017; Arı and Sonmez 2025). Built upon evidence that citizens also differentiate tools by objective (Levin and Musgrave 2025), I argue that instruments carry not only cost-benefit weight but different relational meanings. Our observational design complements the experimental literature by showing how warmth and blame jointly reshape tool-configuration preferences outside the experimental frame. Some tools are read as punitive, relationship-ending signals, while others are interpreted as relationship-managing ones—compliance-seeking leverage deployed within an ongoing engagement. The literature has begun to document this distinction but has not yet specified when and for whom the relational framing applies, nor has it clearly theorized why citizens would combine punitive and cooperative tools simultaneously on a single issue—the question this article takes up.

These findings raise a scope question: if such differentiated reasoning exists, does it charac-

terize the general public, or a specific subset of citizens for whom the issue is especially salient? Krosnick's issue-publics framework provides the relevant boundary condition. Citizens vary in the personal importance they attach to particular policy issues. Those for whom an issue is highly important form an "issue public" marked by greater attitude stability and heightened political engagement on that issue (Krosnick 1989, 1990; Krosnick and Telhami 1995). Issue-bounded conditionality builds on this accessibility-importance tradition: high-importance issues are cognitively structured and chronically accessible, a property that permits blame-driven conditionality to organize preferences within the issue domain without propagating to adjacent issues. Issue publics are not fixed demographic categories but form dynamically through a two-stage process. Groups with latent connections to a political issue constitute potential issue publics until a pivotal circumstance personalizes the issue's political relevance for individual members (Rossiter and Harden 2026). By this logic, the fentanyl crisis—personally salient for millions of affected families and newly politicized through attribution to China—may have activated precisely this kind of domain-specific attentive constituency. The concept matters here not only because it helps identify who engages in differentiated reasoning, but also because it suggests why issue domains themselves may become independent axes of preference formation. When an issue commands sustained personal importance, it can generate a distinct attitudinal space, organized around that domain's specific objectives and instruments, one not reducible to general foreign policy posture. If so, the differentiated reasoning documented in the preceding literature is more likely to appear among those for whom the fentanyl crisis is personally salient than across the mass public uniformly.

Existing work therefore suggests that foreign-policy attitudes are organized across distinct objects of judgment. Citizens who are attentive to a given issue may form views about the bilateral relationship as a whole *and* separate views about the appropriate instrument for handling a particular dispute within that relationship. The literature has not theorized how these two levels interact—nor what conditions enable citizens to hold both punitive and cooperative preferences simultaneously. The following conceptual framework addresses this question.

2.2 Issue-Bounded Conditionality: A Conceptual Framework

This article advances two related claims. The first is structural: foreign-policy attitudes toward a rival state operate at more than one analytical level, and under some crisis conditions these levels can come apart. The second is conditional: this decoupling is not uniform across issues but occurs only where specific scope conditions are met. I call the resulting pattern *issue-bounded conditionality*—the coexistence of domain-specific coercive support with cooperative engagement on the same issue, in a manner that does not generalize across policy domains. The two claims are logically separable and face different counterarguments, so I develop them in turn.

Two levels of opinion. Citizens differentiate by objective, by instrument, and by issue domain. These findings point to two analytically distinct levels of opinion: a *generalized bilateral orientation* toward the rival state and an *issue-level instrument preference* over the specific policy tools endorsed for a bounded problem. Existing research generally treats these levels as reflecting a common underlying orientation, so warmer feelings should produce softer preferences across the board and blame should make both general posture and instrument choice more confrontational. Yet survey evidence already suggests that Americans hold ambivalent orientations toward China (Wick 2014; Chubb 2024; Hu 2025). Parallel evidence from China suggests that similarly mixed orientations toward America are also possible: nearly 46% of Chinese respondents who classified the United States as a “rival” nonetheless expressed a favorable overall attitude toward it (Wick 2014; Wang et al. 2022). The central structural claim here is that these two levels need not move together. Under some conditions, a citizen’s relationship-level orientation may diverge from issue-level instrument preferences—supporting both coercive and cooperative tools within a bounded policy domain. The most obvious counterargument is that such apparent divergence is a measurement artifact: respondents who endorse many items on any list will tend to endorse both sanctions and cooperation, producing a spurious appearance of differentiation. I address this concern empirically through acquiescence controls and the domain-level specificity test that follows, but the structural claim itself stands or falls on whether the two levels carry distinct predictive information once response style is netted out.

Domain-boundedness. The second claim is that this differentiation is not free-floating but tethered to a specific issue. Issue-bounded conditionality is defined structurally as an observable preference configuration rather than a confirmed psychological mechanism. It is distinct from dispositional accounts in which cooperative and militant orientations are stable trait-level differences rooted in personal values (Rathbun et al. 2016), and also distinct from generalized ambivalence toward an adversary: ambivalence describes mixed evaluations at the relationship level, whereas issue-bounded conditionality specifies mixed *tool preferences* organized within a specific issue domain. The counterargument here is that citizens who hold mixed views on fentanyl may hold similarly mixed views on any China-related issue they find salient, so the apparent domain-boundedness is simply a reflection of which issue happens to be in front of them. The domain-level specificity design in Section 4.3 is constructed to adjudicate this: if the S+A pattern is merely a face-valid response to any China scenario, it should appear across all eight probes, not only the fentanyl one.

Why blame conditions warmth. What mechanism could produce this configuration? I treat *bounded instrumental logic*—or, more plainly, leverage—as the working hypothesis. Blame attribution crystallizes a concrete problem and names a responsible actor; warmth provides a relational baseline that respondents are reluctant to abandon. Citizens who hold both at once have a motive to pursue the one tool set that reconciles them: pressure applied *within* rather than *instead of* an ongoing relationship. On this reading, sanctions paired with a demand for Chinese government action are not two contradictory impulses but a single coherent stance—compliance-seeking leverage that preserves the bilateral channel while raising the cost of inaction (**tomz_domestic_2007**; Heinrich, Kobayashi, and Peterson 2017; McLean and Roblyer 2017). Coercive instruments framed as compliance-seeking tools attract more public support than those presented as expressive punishment (Frye 2019). An alternative mechanism is normative rather than instrumental: some respondents may endorse sanctions plus action because both items are cued as the “right” response to an adversary’s transgression, independent of any calculation about compliance. A cross-sectional design cannot adjudicate directly between these readings, since both generate similar patterns of item endorsement. I present descriptive evidence in Section 4.4 that is more consistent with the

leverage reading, and return to the identification problem in the Discussion. For the remainder of the paper, I use *cooperation* and *leverage* interchangeably to refer to this compliance-seeking form of engagement—pressure applied within an ongoing bilateral channel rather than a relational overture rooted in trust.

This decoupling should be most likely under two broad conditions, each specifying when the structural pattern described above is most likely to emerge. The first concerns *problem structure*. The crisis should be domestically salient, responsibility should be traceable to an identifiable foreign actor, and resolution should require that actor’s active participation rather than merely its punishment or exclusion. Where these elements converge, the conditions for mixed configurations are most favorable. When any of them is absent—the issue lacks personal salience, attribution is diffuse, or punishment alone suffices—the foreign-policy dimension either does not activate or does not require cooperative tools, and the framework does not predict mixed configurations. The second concerns *tool structure*. The available coercive instruments should be perceived as limited, reversible, and issue-specific rather than relationship-ending. Sanctions that are seen as bounded can coexist with cooperation in a way that comprehensive measures—a full economic embargo or military confrontation—cannot. When tools are perceived as escalatory or relationship-terminating, citizens are more likely to revert to a unidimensional confrontational posture. Where either condition is absent, standard threat-affect dynamics may dominate. By contrast, a military confrontation over Taiwan—where the available tools are closer to relationship-ending than bounded—would not satisfy the tool-structure condition, and the framework would predict reversion to a unidimensional confrontational posture rather than mixed configurations. The U.S.-China fentanyl case provides a useful empirical setting because it plausibly combines both conditions: high personal salience, attributable foreign actor, compliance-dependent problem structure, and a policy menu that includes bounded coercive instruments (Martínez-Fernández and Harding 2024; Hu, Baldwin, and Seyle 2025; Vangelov et al. 2026).

If this account is correct, it yields two core observable expectations (OEs). The first OE relates to general posture. Warmth and blame attribution should jointly structure foreign-policy ori-

entation, such that blame amplifies rather than eliminates the warmth-posture gradient. Second, if issue-bounded conditionality characterizes issue-level preferences, mixed coercive-cooperative support should remain concentrated in the focal issue domain rather than extend across unrelated ones. In the present case, this implies that support for a sanctions-plus-cooperation configuration should be concentrated on fentanyl-related policy choices. OE2 is the framework’s primary falsifiable implication: IBC does not predict that a particular cooperative item is uniquely effective, only that whatever engagement-compatible configuration emerges should be domain-specific. A finding that blame-driven mixed configurations appear with equal force across military, economic, and cultural domains would disconfirm IBC; a finding that alternative engagement-compatible phrasings produce similar blame-driven shifts within the fentanyl domain would not.

3 Data and Methods

3.1 Sample and Survey Design

I draw on two cross-sectional surveys of U.S. adults fielded by YouGov. The first wave was fielded from March 12 to March 19, 2024 (3,237 interviewed; matched to a final sample of 3,000; reported margin of error approximately $\pm 1.94\%$). The second wave was fielded from April 15 to April 25, 2025 (2,273 interviewed; matched to a final sample of 2,000; reported margin of error approximately $\pm 2.4\%$).

In both waves, respondents were matched to a politically representative modeled frame of U.S. adults built from the American Community Survey (ACS) PUMS, public voter-file records, the 2020 CPS Voting and Registration Supplement, the 2020 National Election Pool (NEP) exit poll, and the 2020 Cooperative Election Study (CES). Matching targets included gender, age, race, and education. Matched cases were weighted using propensity scores via logistic regression using age, gender, race/ethnicity, education, region, and home ownership (home ownership included in 2024 only), grouped into deciles, and post-stratified within deciles. Final weights were post-stratified by presidential vote choice and by a four-way stratification of gender, age, race, and education. All

analyses use these final survey weights.

Question wording for the warmth, blame, posture, and tool-configuration items is identical across the 2024 and 2025 waves. The 2024 wave was fielded in March 2024, before the Trump administration’s fentanyl-related tariffs; the 2025 wave was fielded April 15–25, 2025, after the early-2025 tariff imposition but before the Busan bilateral presidential agreement later that year. Cross-wave differences in raw means may therefore reflect both real attitude change and contextual framing differences arising from the changed political environment. The within-wave replication (Appendix Table A14) confirms that the warmth $\times$ blame pattern holds in each wave independently.

Table 1 reports the weighted composition of both waves across party identification, race/ethnicity, education, gender, age, ideology, and household income. The two samples are demographically similar and track the national adult population on standard benchmarks: roughly 51% female, a quarter under age 30, about a third with a four-year college degree, and a partisan balance close to the national distribution, with a modest shift toward Republican identification in the 2025 wave.

3.2 Key Measures

The two primary predictors are *warmth* toward China, a 7-point favorability measure standardized using pooled-wave moments (*warmth_z*) to ensure comparability across waves, and *blame attribution*, a binary indicator of whether respondents attribute responsibility for the U.S. fentanyl crisis to China (from the multi-select attribution battery). Their relative roles shift across analytical steps. Warmth is the primary predictor of general foreign-policy posture (Step 1), while blame is the primary driver of tool-configuration shifts (Step 2); the interaction between them is the key estimand in both. An alternative single-choice “most responsible” measure is used in robustness analyses.

To capture the multidimensionality of foreign policy preferences, I construct two sets of dependent variables:

Foreign-policy posture index. A four-item index capturing how respondents think America *should* act toward China—aggressive vs. defensive, competitive vs. cooperative, friendly vs. unfriendly, and destructive vs. constructive. Higher values indicate a more dovish/cooperative orien-

Table 1: Sample Demographics: Weighted Proportions (Unweighted N)

Variable	Category	2024	2025
Party ID	Democrat	39.7% (1056)	35.3% (656)
Party ID	Republican	30.4% (845)	34.3% (576)
Party ID	Independent	29.9% (817)	30.4% (565)
Race/Ethnicity	White	65.7% (2019)	65.8% (1368)
Race/Ethnicity	Black	12.5% (381)	12.5% (215)
Race/Ethnicity	Hispanic	11.4% (294)	11.1% (207)
Race/Ethnicity	Asian	3.3% (102)	2.6% (54)
Race/Ethnicity	Other	7.1% (204)	8.0% (156)
College Degree	College+	33.8% (1053)	33.3% (691)
Female	Female	51.3% (1601)	51.3% (1075)
Age Group	18-29	19.4% (513)	19.5% (299)
Age Group	30-44	25.0% (751)	24.6% (527)
Age Group	45-64	32.3% (1050)	33.5% (723)
Age Group	65+	23.3% (686)	22.4% (451)
Ideology	Liberal	29.6% (818)	30.6% (577)
Ideology	Moderate	36.5% (991)	34.6% (655)
Ideology	Conservative	33.9% (958)	34.8% (611)
Household Income	Under \$50k	43.1% (1176)	40.0% (748)
Household Income	\$50k-\$100k	31.1% (834)	31.9% (570)
Household Income	Over \$100k	25.8% (670)	28.1% (516)

tation. The index is standardized as *posture_z* using pooled-wave moments. The four items have Cronbach's $\alpha = 0.87$ (pooled sample), and the index loads cleanly on a first principal component explaining 73% of item variance.

Tool configuration. Preferences for particular foreign policy tools were captured through a checklist where respondents could endorse multiple approaches to the fentanyl crisis. I focus on two items in the foreign-policy domain: support for sanctions or other enforcement actions and support for requiring action by the Chinese government. The phrasing “requiring action by the Chinese government” signals a conditional demand rather than a cooperative overture, but the channel it presupposes is bilateral engagement: one can only “require” action from an actor with whom a working relationship is maintained. Crossing these yields four mutually exclusive categories: *Neither*, *Sanctions* only, *Action* only and *Sanctions+Action* (*S+A*). This decomposition allows us to

distinguish mixed policy configurations from single-tool preferences.

Interpreting the Sanctions + Action configuration. A reasonable reader might object that “requiring action by the Chinese government” is itself a form of demand, not cooperation, and that pairing it with sanctions therefore describes two flavors of pressure rather than a mixed posture. I take this objection seriously because it bears directly on what the S+A category means. My reading is the following. Sanctions alone express rupture: the United States withdraws value from the relationship and accepts the costs of doing so. A demand for Chinese government action expresses something different—it presupposes an ongoing bilateral channel through which compliance can be sought, verified, and rewarded. You cannot “require action” from an actor with whom you have no working relationship; you can only punish them. In this narrower, instrumental sense, endorsing the action item alongside sanctions signals a willingness to keep the channel open while raising the cost of inaction. This is the reading I adopt. Throughout the rest of the paper, *cooperation* refers to this form of conditional engagement—pressure applied *within* rather than *instead of* an ongoing relationship—and I use “cooperation” and “leverage” interchangeably with that narrower meaning. The alternative reading, that the action item simply registers a second dose of pressure, would predict that S+A respondents behave similarly to sanctions-only respondents on measures of engagement with China; the descriptive evidence in Section 4.4 and the placebo comparison in the Discussion speak to that reading. Cognitive interviewing to verify how individual respondents interpreted this item was not conducted and represents a design limitation; the behavioral evidence in Section 4.4 provides the closest available construct validity check from the respondent end.

The core controls include party identification, ideology, education, age, gender, race/ethnicity, political attention, and fentanyl-related exposure indicators. Pooled models include year fixed effects. The use of a multi-item checklist raises the standard concern that acquiescent response style inflates endorsement rates (Bishop 2005); question-order and response-format effects in attitude surveys are well-documented and can further compound checklist-based measurement (Schuman and Presser 1981). To address generalized endorsement propensity, I use endorsement-count controls (main specification excludes the two items that define tool configuration).

3.3 Analytical Approach

The empirical design proceeds in three linked steps. First, I examine whether affective orientation toward China and crisis-specific blame jointly structure general foreign-policy posture. I estimate survey-weighted OLS models (wave-specific and pooled) predicting the posture index to establish the baseline interactive effect of *warmth* $\times$ *blame*:

$$\text{posture_z}_i = \beta_0 + \beta_1 \text{warmth_z}_i + \beta_2 \text{blame_china}_i + \beta_3 (\text{warmth_z}_i \times \text{blame_china}_i) + \Gamma \mathbf{X}_i + \delta_{\text{year}} + \varepsilon_i.$$

The interaction term is the key estimand. Robustness analyses use an alternative blame operationalization (single-choice “most responsible” item) and an extended specification that adds an index of perceptions of Chinese behavior. Both are reported in the Appendix.

The second step examines whether the warmth-by-blame interaction identified in Step 1 is reflected in discrete policy-tool configurations, and constructs the S+A classification that the specificity test requires as input. I estimate multinomial logistic regression models for the four-category tool-configuration outcome (Neither, Sanctions-only, Action-only, S+A), using survey weights as frequency weights, and report predicted probabilities and blame contrasts at $\text{warmth_z} \in \{-1, 0, +1\}$.¹ The S+A category is theoretically central, for it represents the “leveraged cooperation” posture in which respondents simultaneously pursue punitive and cooperative tools. All specifications include *posture_z* (standardized using pooled-wave moments, as described above); accordingly, Step 2 estimates are interpreted as conditional associations in tool composition given overall policy orientation (i.e., within-posture composition effects), not total or causal effects. A post-treatment

1. The multinomial logit assumes independence of irrelevant alternatives (IIA). Appendix Table A18 reports Hausman-McFadden formal tests. IIA is rejected for the Neither ($\chi^2 = 1017.6$, $df = 36$, $p < 0.001$) and S+A ($\chi^2 = 1257.5$, $df = 36$, $p < 0.001$) alternatives, but not for Sanctions-only ($\chi^2 = 19.4$, $p = 0.989$) or Action-only ($\chi^2 = -3.5$, $p = 1.000$). Rejection of IIA implies that substitution across alternatives is not proportional—removing a given option would redistribute its probability mass unevenly across the others—rather than that the model is invalid. I therefore interpret the MNL estimates as conditional associations within the four-alternative menu actually presented to respondents, not as structural preferences that would extrapolate to settings where the option set changes. The coefficient-stability heuristic in Appendix Table A19 confirms that estimates are stable to dropping S+A, supporting use of the MNL results within the observed menu. Substantively, the rejection pattern is consistent with the theoretical framing: S+A and Neither occupy distinct attitudinal regions, whereas Sanctions-only and Action-only behave as closer substitutes.

sensitivity check omitting `posture_z` yields virtually identical blame contrasts ($\Delta < 0.001$ difference in $\Delta_{pp}(S+A)$ at all warmth levels), confirming that within-posture conditioning does not materially distort the estimated blame effect (Appendix Table A16). Party identification is included to address the possibility that partisan dispositional tendencies drive the mixed configuration.

Third, I deploy an eight-probe specificity test to see if the observed pattern is issue-specific or generalized endorsement disposition.² The seven non-fentanyl probes serve as within-design negative cases for the scope conditions specified in Section 2.2: they describe scenarios where the full set of conditions (domestically salient crisis, traceable attribution, compliance-dependent resolution, bounded tools) is not met. Respondents rated how their opinion of China would change (1–7 scale) across eight hypothetical scenarios spanning security, economic, cultural, and public-health domains, presented in randomized order. Restricting the sample to 1,846 S+A and Sanctions-only respondents, I stack the eight scenario responses in long format (one row per respondent-scenario), and estimate a survey-weighted OLS model with respondent-level clustering:

$$\text{probe}_{is} = \alpha_s + \lambda_s \text{is_SA}_i + \Pi_s \mathbf{Z}_i + u_i + e_{is},$$

where s indexes scenarios, $\text{is_SA}_i = 1$ for S+A (0 for Sanctions-only), and $\mathbf{Z}_i$ denotes controls. This specification yields scenario-specific contrasts $\Delta_s = \mathbb{E}[Y \mid S+A] - \mathbb{E}[Y \mid S\text{-only}]$. I summarize non-target scenarios with a composite boundary mean and test concentration on the fentanyl target probe using the omnibus divergence:

$$\Delta_d - \frac{1}{7} \sum_{x \in \{a,b,c,e,f,g,h\}} \Delta_x.$$

I also report two one-sided tests (TOST) as an equivalence test to confirm that non-target probe effects are negligibly small (**lakens_equivalence_2018**; Lakens 2017).³ The design can rule out

2. I refer to this as a “specificity test” rather than a “falsification test” because its target is concentration rather than refutation. A null result on the seven boundary probes would falsify a generalized-endorsement account of the S+A pattern, but the design itself is a test of whether the effect is concentrated in the fentanyl domain, not a single prediction that could be rejected in the strict Popperian sense.

3. The equivalence bound is set at $\varepsilon = 0.25 \times SD_{\text{boundary}}$, consistent with the convention that a “small effect”

generalized acquiescence but not domain-specific social desirability.

Listwise deletion reduces the analytic sample from 5,000 to 4,234 respondents (15.3%).⁴ All results are associational. The design uses repeated cross-sections, not panel data. Model-based standard errors in the multinomial specification may understate uncertainty relative to design-based alternatives.

4 Results

4.1 Warmth, Blame, and Foreign Policy Posture

Table 2 reports survey-weighted OLS estimates from six model specifications. The core specification (columns 1–3) estimates baseline models for 2024, 2025, and the pooled sample. Across all three, warmth toward China is the strongest predictor: a one-standard-deviation increase in warmth is associated with a 0.35-unit increase in the posture index ($p < 0.001$, pooled).

The key theoretical quantity is the warmth $\times$ blame interaction. The estimated interaction coefficient is $b_{\text{interaction}} = 0.130$ ($p < 0.001$, pooled): respondents who attribute greater responsibility to China for the fentanyl crisis exhibit a stronger warmth–posture association. In substantive terms, the estimated warmth slope is 0.35 when blame = 0 and 0.48 when blame = 1, a 37% difference in the warmth–posture gradient. Figure 1 illustrates this pattern, showing interaction plots faceted by survey year. Blame attribution alone, by contrast, has a small direct effect on posture that is not

corresponds to approximately 0.2 SD (**lakens_equivalence_2018**). Appendix Table A12 reports sensitivity across four bounds (0.10, 0.15, 0.20, $0.25 \times SD$); the composite boundary estimate is equivalently small at all four choices. The stacked design pools eight probe responses for each of $n = 1,846$ S+A and S-only respondents (effective $n \approx 14,768$ probe-respondent observations with respondent-level clustering), which yields ample power to detect deviations from equivalence at the chosen bounds; failure to reject the equivalence null at $\varepsilon = 0.10 \times SD$ would be informative even at this stringent threshold, and the test in fact rejects at every bound considered.

4. Dropped respondents are younger (mean age 44 vs. 51), warmer toward China (3.5 vs. 2.9), and less likely to blame China (37% vs. 51%). Item nonresponse on politically sensitive items is not random and can selectively mask the views of less politically engaged respondents (Berinsky 2002). If anything, this pattern biases the estimated warmth $\times$ blame interaction downward, since warmer and less-blaming respondents are under-represented after deletion. A sensitivity analysis omitting ideology—the primary source of missingness—recovers additional cases with substantively consistent estimates. A multiple imputation robustness check (10 imputations, Predictive Mean Matching) recovers the full sample of 5,000 respondents and yields a warmth $\times$ blame coefficient of $b = 0.123$ ($SE = 0.026$), an attenuation of approximately 8% relative to the complete-case estimate. The fraction of missing information is 0.001, indicating that missingness contributes negligible uncertainty to the pooled estimate (Appendix Table A20).

statistically significant at conventional levels ($b = -0.050$, $p = 0.096$). This asymmetry—blame amplifies warmth rather than independently shifting policy preferences—is consistent with the first OE: responsibility attributions operate as a moderator of the warmth–posture gradient rather than as an independent driver of posture. Ideological self-placement is the only other consistently significant predictor ($b = -0.100$, $p < 0.001$, pooled), though smaller in magnitude than warmth.

The interaction is robust across specifications. It holds under an alternative blame operationalization using the single-choice “most responsible” item (Table 2), when a behavioral engagement index is added as a further control (Table 2, columns 4–6), and in ordinal logistic regression models treating individual posture items as ordered outcomes (Appendix Table A1).




Figure 1: Predicted foreign policy posture by warmth toward China and blame attribution, faceted by survey year. Lines represent survey-weighted OLS predictions with 95% confidence intervals. The steeper slope under high blame illustrates the amplification interaction.

Table 2: Survey-Weighted OLS: Foreign-Policy Posture Index on Warmth and Blame

	A:2024	A:2025	A:Pool	B:2024	B:2025	B:Pool
(Intercept)	0.71*** (0.10)	0.38** (0.14)	0.54*** (0.08)	0.54*** (0.10)	0.21 (0.14)	0.38*** (0.08)
warmth_z	0.39*** (0.03)	0.30*** (0.04)	0.35*** (0.03)	0.21*** (0.03)	0.19*** (0.04)	0.20*** (0.03)
blame_china	-0.05 (0.04)	-0.07 (0.05)	-0.05 (0.03)	-0.01 (0.04)	-0.03 (0.05)	-0.02 (0.03)
party3Republican	-0.11 (0.06)	-0.09 (0.08)	-0.10* (0.05)	-0.07 (0.06)	-0.08 (0.08)	-0.07 (0.05)
party3Independent	-0.04 (0.05)	-0.06 (0.06)	-0.05 (0.04)	0.01 (0.05)	-0.02 (0.06)	-0.00 (0.04)
ideo5_clean	-0.13*** (0.02)	-0.06 (0.03)	-0.10*** (0.02)	-0.12*** (0.02)	-0.04 (0.03)	-0.09*** (0.02)
educ_college	-0.04 (0.04)	0.01 (0.05)	-0.02 (0.03)	-0.03 (0.04)	0.02 (0.04)	-0.01 (0.03)
age	-0.00** (0.00)	-0.00* (0.00)	-0.00*** (0.00)	-0.00* (0.00)	-0.00* (0.00)	-0.00** (0.00)
female	0.01 (0.04)	0.07 (0.05)	0.04 (0.03)	0.03 (0.04)	0.06 (0.04)	0.04 (0.03)
race_ethBlack	0.01 (0.06)	-0.11 (0.08)	-0.04 (0.05)	0.00 (0.06)	-0.12 (0.08)	-0.05 (0.05)
race_ethHispanic	0.00 (0.08)	-0.08 (0.08)	-0.03 (0.06)	0.01 (0.08)	-0.07 (0.08)	-0.02 (0.06)
race_ethAsian	-0.34** (0.11)	-0.24* (0.10)	-0.30*** (0.08)	-0.27* (0.11)	-0.18 (0.11)	-0.24** (0.08)
race_ethOther	-0.13 (0.07)	0.07 (0.08)	-0.04 (0.05)	-0.10 (0.07)	0.08 (0.08)	-0.02 (0.05)
newsint_attn	-0.02 (0.02)	0.05 (0.03)	0.01 (0.02)	-0.01 (0.02)	0.06* (0.03)	0.02 (0.02)
exposure_index	0.02 (0.03)	-0.02 (0.04)	0.00 (0.03)	0.01 (0.03)	-0.03 (0.04)	-0.01 (0.02)
warmth_z:blame_china	0.10* (0.05)	0.18*** (0.05)	0.13*** (0.03)	0.10* (0.04)	0.15** (0.05)	0.12*** (0.03)
factor(year)2025			0.11*** (0.03)			0.11*** (0.03)
behavior_index_z				0.30*** (0.03)	0.24*** (0.03)	0.28*** (0.02)
Deviance	1913.21	1203.12	3141.58	1783.39	1148.17	2956.09
Dispersion	0.75	0.71	0.74	0.70	0.68	0.70
Num. obs.	2545	1689	4234	2545	1689	4234

*** $p < 0.001$; ** $p < 0.01$; * $p < 0.05$

	A:2024	A:2025	A:Pool	B:2024	B:2025	B:Pool
(Intercept)	0.70*** (0.11)	0.39** (0.14)	0.54*** (0.09)	0.55*** (0.11)	0.22 (0.14)	0.38*** (0.09)
warmth_z	0.42*** (0.03)	0.37*** (0.03)	0.40*** (0.02)	0.25*** (0.03)	0.23*** (0.03)	0.24*** (0.02)
blame_china_alt	-0.06	-0.09	-0.06	-0.09	-0.09	-0.08

4.2 Issue-Level Tool Configuration

Having established the warmth–blame interaction at the posture level, I now examine whether this pattern is reflected in discrete tool configurations. Before turning to the multinomial results, a preliminary variance decomposition confirms that the two dependent variables capture largely distinct aspects of foreign-policy preferences. Tool configuration explains only 2.5% of the weighted variance in general posture ($\eta^2 = 0.025$, $p < 0.001$). S+A and Sanctions-only respondents occupy similar positions on the general posture spectrum (weighted means of -0.24 and -0.14 respectively), yet they differ sharply in tool composition.

Table 3 reports predicted probabilities and blame-effect contrasts at one standard deviation below and above the warmth mean. Blame attribution is associated with a substantially higher probability of S+A classification at all warmth levels. Among cold respondents (warmth = -1 SD), the estimated $\hat{P}(\text{S+A})$ is 0.069 without blame and 0.363 with blame ($\Delta = +0.294$, $p < 0.001$). Among warm respondents ($+1$ SD), the corresponding figures are 0.044 and 0.350 ($\Delta = +0.306$, $p < 0.001$). By contrast, the blame effect on the Sanctions-only category is small and non-significant ($\Delta \approx 0$), consistent with blame channeling respondents into the combined S+A posture rather than a purely punitive stance.⁵

Figure 2 displays the predicted probability of each tool configuration across *warmth* $\times$ *blame* quadrants, making the selectivity of the blame effect visually apparent.

A potential concern is that the S+A pattern reflects acquiescent response style—a general tendency to endorse survey items—rather than genuine policy differentiation (Krosnick 1991). I address this by controlling for an endorsement-count measure that attenuates the S+A blame effect by approximately 38% (from $\Delta = 0.306$ to $\Delta = 0.191$ among warm respondents). Such a substantial reduction indicates that acquiescent response style accounts for a meaningful share of the raw effect.

5. The blame *main effect* on $P(\text{S+A})$ is robust to design-based inference (Appendix Table A5). The warmth $\times$ blame *interaction* on the S+A category, however, is not statistically distinguishable from zero under design-based standard errors ($b = 0.134$, $SE = 0.114$, $t = 1.18$, $p = 0.240$, from a survey-weighted binary contrast on $P(\text{S+A})$), in contrast to the model-based MNL estimate. The point estimate of the warmth–blame interaction in tool composition is therefore best read as suggestive rather than precisely estimated; the substantively important quantity for the issue-bounded conditionality argument is the large blame main effect on S+A, which is precise under both inference approaches.




Figure 2: Predicted probability of each tool-configuration category by warmth $\times$ blame quadrant, from survey-weighted multinomial logit. Error bars show 95% confidence intervals. Blame shifts respondents into S+A (sanctions plus cooperation) rather than Sanctions-only.

Table 3: Predicted Probabilities and Blame-Effect Contrasts: Multinomial Logit Estimates

Estimand	Estimate	SE	95% CI
P_SA_w-1_b0	0.0688	0.0112	[0.0469, 0.0906]
P_Sanctiononly_w-1_b0	0.1907	0.0201	[0.1513, 0.2301]
P_SA_w-1_b1	0.3632	0.0287	[0.3070, 0.4194]
P_Sanctiononly_w-1_b1	0.2229	0.0212	[0.1813, 0.2645]
P_SA_w+1_b0	0.0437	0.0073	[0.0293, 0.0581]
P_Sanctiononly_w+1_b0	0.2097	0.0209	[0.1688, 0.2507]
P_SA_w+1_b1	0.3495	0.0321	[0.2866, 0.4124]
P_Sanctiononly_w+1_b1	0.1930	0.0229	[0.1482, 0.2379]
Dpp_SA_w-1	0.2944	0.0248	[0.2458, 0.3430]
Dpp_Sonly_w-1	0.0322	0.0203	[-0.0075, 0.0719]
Dpp_SA_w+1	0.3058	0.0289	[0.2492, 0.3624]
Dpp_Sonly_w+1	-0.0167	0.0211	[-0.0581, 0.0247]

Notes. Survey-weighted multinomial logit. Predicted probabilities evaluated at warmth = ± 1 SD; contrasts (Δ_{pp}) are predicted-probability differences between blame = 1 and blame = 0 at the indicated warmth level. Standard errors are Hessian-based (model-based). Reference category: Neither.

Nevertheless, the qualitative pattern—blame increases S+A, not Sanctions-only—remains intact, and the specificity test in Section 4.3 provides further evidence against a pure response-style explanation (Appendix Table A3). A placebo test replacing the dependent variable with a sanctions $\times$ international-cooperation composite (theoretically unrelated to specific U.S.-China collaboration) does not support tool-specificity. I return to this limitation in Section 5.

4.3 Testing Issue-Bounded Conditionality Across Scenarios

The preceding analyses show that blame shifts respondents into the S+A configuration. However, does this classification reflect issue-bounded conditionality, or a general willingness to cooperate with China? The specificity test addresses this by comparing the S+A-versus-S-only contrast across eight scenarios. If the contrast appears only on the fentanyl target probe and not on the seven boundary probes, the classification captures something specific to that issue rather than a diffuse disposition.

Table 4 maps each scenario to the scope conditions developed in Section 2.2. The assignment of scope conditions to probes is theory-derived: the two conditions (problem structure and tool structure) were specified in Section 2.2 prior to any analysis of the probe battery, and the probe-to-condition mapping follows directly from the theoretical definitions. The probes were not assigned to conditions based on the pattern of the results; the expected direction (≈ 0 vs. > 0) was fixed before estimation. The target probe (probe *d*, *Fentanyl cooperation*) is the only scenario in which both conditions are fully satisfied. The three security probes (*a–c*) violate the *tool-structure* condition: military instruments are relationship-ending rather than bounded, so the framework predicts reversion to a unidimensional confrontational posture. The four non-military probes (*e–h*) violate the *problem-structure* condition: they describe steady-state policy questions without a salient crisis, traceable blame, or compliance-dependent resolution. In all seven boundary cases, the framework predicts no amplified S+A-minus-S-only contrast.

Table 5 reports the core results, where each entry is the adjusted mean difference in opinion-change scores between S+A and Sanctions-only respondents for that probe. The target contrast is

Table 4: Eight-Probe Battery: Domain, Scope Conditions, and Expected Direction

Probe	Scenario	Domain	Scope condition violated	Expected Δ
<i>a</i>	Taiwan military action	Security	Tool structure	≈ 0
<i>b</i>	U.S.-China military collision	Security	Tool structure	≈ 0
<i>c</i>	Third-party military collision	Security	Tool structure	≈ 0
<i>d</i>	Fentanyl cooperation	Health	None (target)	> 0
<i>e</i>	Resuming adoptions from China	Cultural	Problem structure	≈ 0
<i>f</i>	Halting electronics sales	Economic	Problem structure	≈ 0
<i>g</i>	Chinese student scholarships	Cultural	Problem structure	≈ 0
<i>h</i>	Pandas returning to U.S. zoos	Cultural	Problem structure	≈ 0

$\Delta_d = +0.282$ ($p = 0.002$). In substantive terms, S+A respondents score 0.28 scale points higher than S-only respondents on the fentanyl-cooperation probe (on a 1–7 scale), after controlling for demographics, ideology, political attention, and exposure. None of the seven boundary probes achieves significance individually, and the composite boundary mean is small and non-significant ($\bar{\Delta}_{\text{boundary}} = +0.042$, $p = 0.337$).

The omnibus divergence—the difference between the target contrast and the composite boundary mean—is $+0.240$ ($p = 0.008$), consistent with a concentration of the S+A effect on the fentanyl cooperation dimension. A TOST equivalence test confirms that the composite boundary mean is negligibly small ($p = 0.009$). It indicates that S+A respondents do not differ meaningfully from S-only respondents on non-fentanyl scenarios. Figure 3 presents the results visually, contrasting the S+A–versus–S-only adjusted means for each probe. Probe *d* is the only probe with a significant and substantively large contrast.

Table 5: Issue-Specificity Test: Adjusted Mean Differences in Opinion-Change Scores (S+A minus S-only)

Contrast	Estimate	SE	CI	p_value
Delta_d (target: Fentanyl cooperation)	0.2817	0.0913	[0.1027, 0.4607]	0.002**
Composite boundary mean(Delta_others)	0.0417	0.0434	[-0.0435, 0.1269]	0.337
Omnibus divergence (Delta_d - composite)	0.2400	0.0910	[0.0616, 0.4184]	0.008**
TOST equivalence (eps=0.145)	0.0417	0.0434	+/- 0.145	0.009**

Notes. Stacked OLS with respondent clustering. Each row is Δ_k = adjusted mean difference in opinion-change score between S+A and Sanctions-only respondents for probe *k*, on a 1–7 scale. Controls: demographics, ideology, political attention, and exposure index. * $p < 0.05$; ** $p < 0.01$; *** $p < 0.001$.




Figure 3: Adjusted mean differences in opinion-change scores (S+A minus S-only, in 1–7 scale-point units) for eight scenario probes. Probe *d* (fentanyl cooperation) is the target; probes *a–c* and *e–h* are boundary probes. Points show estimated Δ with 95% confidence intervals. The dashed vertical line marks zero. Only the target probe is significant, consistent with issue-bounded conditionality.

These results are robust across multiple specifications. Pairwise divergence tests—comparing Δ_d against each boundary probe individually, with Benjamini–Hochberg false discovery rate correction (Benjamini and Hochberg 1995)—confirm that the target probe is significantly larger than two of seven boundary probes at the BH-FDR adjusted $\alpha = 0.05$ level (four of seven at the unadjusted level), and directionally larger than all seven (Appendix Table A8). Alternative specifications using zero-imputed endorsement counts, continuous $\hat{P}(S+A)$ predicted probabilities in place of the binary classifier, and the full twelve-item endorsement count yield estimates consistent in sign and magnitude (Appendix Tables A9–A11). In sum, the concentration of the S+A contrast on the fentanyl probe, with near-zero boundary effects, is consistent with issue-bounded conditionality rather than a generalized cooperative disposition.

4.4 Interpreting the S+A Configuration: Descriptive Mechanism Evidence

Section 2.2 advanced *bounded instrumental logic*—leverage—as the working hypothesis for why blame conditions warmth, and flagged a normative alternative in which sanctions plus action simply encode the “right” response to adversary transgression. A cross-sectional design cannot adjudicate directly between these mechanisms, since both predict elevated joint endorsement of the two items. What the survey can do is report, descriptively, whether S+A respondents view the bilateral channel in ways consistent with the leverage reading—whether they treat China as an actor whose behavior they are trying to move, rather than as an opponent to be punished expressively. I use four items from the instrument to probe this: whether respondents name the Chinese government as primarily responsible for fixing the problem; whether they endorse improved international cooperation as part of the solution; whether they endorse health and justice collaborations; and where they place China on a seven-point competitive–cooperative scale. Table 6 reports survey-weighted comparisons between S+A and S-only respondents on these four items, pooled across the 2024 and 2025 waves.⁶

Three features of the comparison support the leverage reading. First, S+A respondents are more

6. The four items are available in both waves (PAXS0003 and PAXS0004); the seven-point competitive–cooperative scale is labeled q14_b in the 2024 instrument and q13_b in the 2025 instrument but uses identical wording and response options.

Table 6: Descriptive Mechanism Comparisons: S+A vs. Sanctions-only Respondents

Item	S+A	S-only	Difference (95% CI)
% naming Chinese government as primarily responsible for fixing (q9)	10.0%	2.8%	+7.2 pp (5.0, 9.4)
% endorsing improved international cooperation (q10_1)	77.2%	38.1%	+39.0 pp (34.8, 43.3)
% endorsing health and justice collaborations (q10_5)	45.3%	29.1%	+16.3 pp (11.8, 20.7)
Mean competitive-cooperative rating, 1–7 (q13_b / q14_b)	2.04	2.55	−0.51 (−0.66, −0.36)

Notes. Pooled sample of S+A and Sanctions-only respondents across the 2024 and 2025 waves ($n = 2,071$; 1,083 S+A, 988 S-only). Entries are survey-weighted proportions (rows 1–3) and a weighted mean (row 4). Confidence intervals are based on linearization standard errors. Positive differences favor the S+A group; a negative difference on row 4 indicates that S+A respondents place China closer to the “competitive” end of the 1–7 scale than S-only respondents.

than three times as likely as S-only respondents to name the Chinese government as the group most responsible for *fixing* the problem (10.0% vs. 2.8%, a 7.2-point gap). This is the strongest single indication that “requiring action by China” reads, for the S+A group, as the *remedy* they are trying to extract rather than as a second punitive demand layered on top of sanctions: the Chinese state is cast as the agent whose compliance resolves the crisis, not merely as the agent who caused it. Second, S+A respondents endorse the international-cooperation and health-collaboration items at substantially higher rates—77.2% vs. 38.1% for improved international cooperation, and 45.3% vs. 29.1% for health and justice collaborations. Their general orientation toward the fentanyl problem is visibly collaborative rather than punitive; in their adjacent answers they look like people who think the solution runs through engagement. Third, this collaborative orientation is *not* rooted in a benign view of the bilateral relationship. On the seven-point competitive-cooperative scale, S+A respondents place China at 2.04, about half a point *more* competitive than the 2.55 mean among S-only respondents. The combination is exactly what a leverage logic would predict: S+A respondents are realist—even a touch more hard-eyed than S-only respondents—about China’s current behavior, yet they endorse engagement-shaped instruments because they see compliance as the objective rather than rupture. A group that already believed the relationship was amicable

would have little reason to pair sanctions with demands in the first place.

I emphasize that these patterns are descriptive, not causal. They do not identify a mechanism—they characterize the orientation of the S+A group on adjacent items. A respondent who answers the four items this way could still be motivated by normative considerations rather than instrumental ones; the items do not directly tap the citizen’s reasoning process. What the comparison does establish is whether the mechanism story is *coherent*: whether S+A respondents look, in their other answers, like people pursuing leverage or like people cheering for punishment. The Discussion returns to what the cross-sectional design can and cannot settle and sketches an experimental design that could adjudicate the question directly.

5 Discussion and Conclusion

General foreign-policy posture and issue-level tool preferences toward China operate as partially independent levels of opinion, and crisis-specific blame is associated with a reorganization of the relationship between them. Blame for the fentanyl crisis does not push warm respondents toward confrontation. It is associated with a distinctive mixed configuration in which coercive and cooperative tools coexist within a bounded policy domain. This two-level structure challenges the assumption, implicit in much of the image-theory and threat-perception literature, that affective orientation toward a rival state determines policy preferences in a unidimensional fashion. When a domestic crisis is attributed to a foreign actor, citizens appear to differentiate not only between policy objectives, as the “pretty prudent public” tradition has long argued, but also between instruments. The present results are more consistent with a reading in which some coercive tools are compatible with ongoing engagement than with one in which they signal relationship rupture.

Two features of the results warrant particular interpretive caution, and both acquire a different weight once the S+A configuration is read as compliance-seeking leverage rather than as cooperation in the relational-trust sense developed in Section 2.2. First, a placebo test using an alternative tool pairing (sanctions combined with *improved international cooperation*, in place of

China-specific fentanyl cooperation) produced a blame-driven shift of comparable magnitude. A formal comparison of the two blame contrasts— $\Delta_{pp}(S+A)$ from the main specification against $\Delta_{pp}(\text{Sanctions}+\text{IntlCoop})$ from the placebo—yields a difference of -0.029 ($SE = 0.095$, $t = -0.31$, $p = 0.758$; Appendix Table A17). The two contrasts are statistically indistinguishable. Read narrowly as a failure of tool-level specificity, this finding would be corrosive: it would suggest that any cooperative-sounding item pairs similarly well with sanctions, and the S+A configuration might then reflect a generic “sanctions plus something nice” response pattern. Read through the leverage framework, however, the placebo is considerably less damaging. Both pairings are engagement-compatible pressure packages: both couple a coercive instrument with a cooperative one, and both presuppose an ongoing channel through which pressure can be applied and compliance extracted. That they move in parallel with blame is consistent with the idea that blame activates a *leverage orientation* that does not discriminate finely between specific cooperative items. What the placebo cannot do, in either reading, is dislodge the domain-level concentration documented in Section 4.3: the eight-probe panorama showed that the S+A contrast is confined to the fentanyl cooperation scenario and is negligible across military, economic, and cultural domains. The framework’s central claim is domain-concentration, not tool-specificity, and that claim survives the placebo.

Second, controlling for acquiescent response style attenuates the S+A blame effect by approximately 38%. The uncorrected estimates should therefore be treated as an upper bound and the acquiescence-adjusted estimates as a conservative floor. Even at the floor, the qualitative pattern persists: blame is associated with an increase in the S+A configuration, not with an increase in Sanctions-only. Under the leverage reading, what the acquiescence correction removes is diffuse willingness to agree with surveyed items; what remains is the part of the S+A pattern that cannot be explained by endorsement style alone, and it is this residual—not the raw coefficient—that carries the framework’s claim. If this pattern generalizes beyond the present sample, it also carries a methodological implication for the broader public-opinion literature: survey-based studies of attitudes toward China that do not account for endorsement tendency may overstate the prevalence of

mixed or cooperative preference configurations.

The evidence supports the claim that warmth and blame jointly structure both posture and tool composition, and that the resulting mixed configuration is domain-specific. It does not establish causality. The design uses repeated cross-sections, and the observed associations may reflect unmeasured confounders or reverse causation. The specificity test reduces the plausibility of a generalized acquiescence account but cannot distinguish bounded instrumental logic from domain-specific social desirability, the possibility that respondents endorse fentanyl cooperation because prevailing media frames make it the expected complement to sanctions, not because of a genuinely differentiated preference structure. The inclusion of the standardized posture index as a control variable conditions on a post-treatment covariate, meaning that the tool-composition estimates (Step 2) capture within-posture effects over total effects. The framework is tested on a single bilateral case and across two waves bracketing a substantial change in the political environment. The 2025 wave was fielded after the Trump administration’s fentanyl-related tariffs but before the later-2025 Busan bilateral agreement, and the warmth $\times$ blame interaction is roughly twice as large in 2025 ($b_{2025} = 0.185$, $p < 0.001$) as in 2024 ($b_{2024} = 0.096$, $p = 0.037$; Appendix Table A14). The triple-interaction test is underpowered ($p = 0.332$), so this between-wave amplification cannot be statistically attributed to the post-tariff context. It is descriptively consistent with the possibility that elite framing tying punitive instruments to counternarcotics cooperation has heightened the conditional pattern documented here, but the present design cannot rule out other political shifts between March 2024 and April 2025 that may have produced a similar effect. Whether issue-bounded conditionality generalizes to other crisis domains or other bilateral relationships—and how it co-varies with elite framing more broadly—remains an open question. Standard errors are model-based rather than design-based, which may understate uncertainty if the assumed variance structure is misspecified.

Public opinion on China is more internally differentiated than aggregate favorability polls indicate (Li 2021), with three implications for scholars and practitioners. The domestic political space for issue-specific engagement may be wider than hawk–dove framings imply. A public hostile

toward China will not necessarily reject cooperative initiatives; what appears to matter is whether cooperation is tied to a concrete, accountable objective—stemming the fentanyl supply—rather than framed as a general diplomatic concession. That latitude is sharply bounded, though: support for cooperation does not transfer across issue areas. Extrapolating from public willingness to cooperate on fentanyl to an appetite for broader engagement on trade, technology, or military confidence-building would misread the data. Blame attribution does not function as a simple barrier to cooperation but as a structuring condition: citizens who blame China most intensely are also, conditional on warmth, the most likely to endorse the mixed sanctions-and-cooperation configuration. This is consistent with the logic of issue publics, in which crisis salience activates a subgroup with distinctive, issue-bounded preferences. “Tough on China” rhetoric and issue-specific cooperation are not zero-sum in the public mind. The former may help make the latter politically acceptable when cooperation is tied to a concrete and accountable objective.

Several extensions would strengthen the framework. Panel data tracking the same respondents over time would clarify whether issue-bounded conditionality is a stable preference configuration or a transient response to crisis salience. Experimental manipulation of blame attribution—randomly assigning respondents to receive information attributing the fentanyl crisis to China, to domestic actors, or to diffuse causes—would permit causal identification of the blame mechanism that the present cross-sectional design can only suggest, and could also adjudicate between the leverage and normative readings of the S+A configuration by measuring whether exogenous shifts in blame move cooperative-item endorsement in the direction leverage predicts. Testing the framework on other bilateral crises that plausibly satisfy the scope conditions, such as U.S.-China cooperation on climate change or technology governance, would establish whether issue-bounded conditionality is a general feature of public opinion under crisis attribution or a phenomenon unique to the case of synthetic opioids. Finally, on the estimation side, hierarchical or multilevel specifications that allow random intercepts for respondents and for scenario probes would be a natural methodological extension of the stacked-OLS design used here, particularly for analyses that aim to partition variance in S+A preferences across individuals, issues, and waves.

References

- Ari, Barış, and Burak Sonmez. 2025. "Human rights violations and public support for sanctions." *Journal of Peace Research* 62 (1): 68–84. <https://doi.org/10.1177/00223433231201450>.
- Barr, Andy. 2025. *H.R.747 - Stop Chinese Fentanyl Act of 2025*, January. <https://www.congress.gov/bill/119th-congress/house-bill/747/text>.
- Benjamini, Yoav, and Yosef Hochberg. 1995. "Controlling the False Discovery Rate: A Practical and Powerful Approach to Multiple Testing." _Eprint: <https://rss.onlinelibrary.wiley.com/doi/pdf/10.1111/j.2517-6161.1995.tb02031.x>, *Journal of the Royal Statistical Society: Series B (Methodological)* 57 (1): 289–300. <https://doi.org/10.1111/j.2517-6161.1995.tb02031.x>.
- Berinsky, Adam J. 2002. "Silent Voices: Social Welfare Policy Opinions and Political Equality in America." *American Journal of Political Science* 46 (2): 276–287. <https://doi.org/10.2307/3088377>.
- Bishop, George F. 2005. *The Illusion of Public Opinion: Fact and Artifact in American Public Opinion Polls*. Lanham, MD: Rowman & Littlefield.
- Bullock, John G. 2011. "Elite Influence on Public Opinion in an Informed Electorate." *American Political Science Review* 105 (3): 496–515. <https://doi.org/10.1017/S0003055411000165>.
- Carnegie, Allison, and Nikhar Gaikwad. 2022. "Public Opinion on Geopolitics and Trade: Theory and Evidence." *World Politics* 74 (2): 167–204. <https://doi.org/10.1017/S0043887121000265>.
- Chubb, Andrew. 2024. *Chinese Nationalism, PRC Resolve, and Crisis Escalation: Views from Indo-Pacific Experts*, August. <https://asiasociety.org/policy-institute/chinese-nationalism-prc-resolve-and-crisis-escalation-views-indo-pacific-experts>.
- Frye, Timothy. 2019. "Economic Sanctions and Public Opinion: Survey Experiments From Russia." *Comparative Political Studies* 52 (7): 967–994. <https://doi.org/10.1177/0010414018806530>.

- Gries, Peter Hays. 2014. ““Red China” and the “Yellow Peril”: How Ideology Divides Americans over China.” *Journal of East Asian Studies* 14 (3): 317–346. <https://doi.org/10.1017/S1598240800005518>.
- Gries, Peter Hays, and H. Michael Crowson. 2010. “Political Orientation, Party Affiliation, and American Attitudes Towards China.” *Journal of Chinese Political Science* 15 (3): 219–244. <https://doi.org/10.1007/s11366-010-9115-1>.
- Guisinger, Alexandra, and Elizabeth N. Saunders. 2017. “Mapping the Boundaries of Elite Cues: How Elites Shape Mass Opinion across International Issues.” *International Studies Quarterly* 61 (2): 425–441. <https://doi.org/10.1093/isq/sqx022>.
- Heinrich, Tobias, Yoshiharu Kobayashi, and Timothy M. Peterson. 2017. “Sanction Consequences and Citizen Support: A Survey Experiment.” *International Studies Quarterly* 61 (1): 98–106. <https://doi.org/10.1093/isq/sqw019>.
- Herrmann, Richard K. 2017. “How Attachments to the Nation Shape Beliefs About the World: A Theory of Motivated Reasoning.” *International Organization* 71 (S1): S61–S84. <https://doi.org/10.1017/S0020818316000382>.
- Herrmann, Richard K., Philip E. Tetlock, and Penny S. Visser. 1999. “Mass Public Decisions to Go to War: A Cognitive-Interactionist Framework.” *American Political Science Review* 93 (3): 553–573. <https://doi.org/10.2307/2585574>.
- Herrmann, Richard K., James F. Voss, Tonya Y. E. Schooler, and Joseph Ciarrochi. 1997. “Images in International Relations: An Experimental Test of Cognitive Schemata.” *International Studies Quarterly* 41 (3): 403–433. <https://doi.org/10.1111/0020-8833.00050>.
- Holsti, Ole R. 2004. *Public Opinion and American Foreign Policy*. Revised Edition. Ann Arbor, MI: University of Michigan Press.

- Hu, Junyang. 2025. *Calibrating Engagement: American Public Opinion of China in Transition (2024-2025)*. Denver, CO, December. <https://paxsapiens.org/calibrating-engagement-american-public-opinion-of-china-in-transition-2024-2025/>.
- Hu, Junyang, Daniel Baldwin, and Conor Seyle. 2025. *The Blame Game Revisited: Evolving U.S. Perspectives on Opioids, Responsibility, and China*. Technical report. Denver, CO: PAX sapiens Foundation, December. <https://paxsapiens.org/the-blame-game-revisited-evolving-u-s-perspectives-on-opioids-responsibility-and-china/>.
- Huang, Christine, Laura Silver, and Laura Clancy. 2025. *Negative Views of China Have Softened Slightly Among Americans*, April. <https://www.pewresearch.org/global/2025/04/17/negative-views-of-china-have-softened-slightly-among-americans/>.
- Huddy, Leonie, Stanley Feldman, Charles Taber, and Gallya Lahav. 2005. “Threat, Anxiety, and Support of Antiterrorism Policies.” *American Journal of Political Science* 49 (3): 593–608. <https://doi.org/10.1111/j.1540-5907.2005.00144.x>.
- Hurwitz, Jon, and Mark Peffley. 1987. “How Are Foreign Policy Attitudes Structured? A Hierarchical Model.” *The American Political Science Review* 81 (4): 1099–1120. <https://doi.org/10.2307/1962580>.
- Jentleson, Bruce W. 1992. “The Pretty Prudent Public: Post Post-Vietnam American Opinion on the Use of Military Force.” *International Studies Quarterly* 36 (1): 49–74. <https://doi.org/10.2307/2600916>.
- Jentleson, Bruce W., and Rebecca L. Britton. 1998. “Still Pretty Prudent: Post-Cold War American Public Opinion on the Use of Military Force.” *Journal of Conflict Resolution* 42 (4): 395–417. <https://doi.org/10.1177/0022002798042004001>.

- Jin, Yongai, Shawn Dorius, and Yu Xie. 2021. "Americans' Attitudes toward the US–China Trade War." *Journal of Contemporary China* 31 (133): 17–37. <https://doi.org/10.1080/10670564.2021.1926089>.
- Kafura, Craig. 2025. *Americans Reverse Course on US-China Competition*, October. <https://globalaffairs.org/research/public-opinion-survey/americans-reverse-course-us-china-competition>.
- Kertzer, Joshua D. 2023. "Public Opinion about Foreign Policy." In *The Oxford Handbook of Political Psychology*, edited by Leonie Huddy, David O. Sears, Jack S. Levy, and Jennifer Jerit, 447–485. Oxford University Press. <https://doi.org/10.1093/oxfordhb/9780197541302.013.12>.
- Kertzer, Joshua D., and Thomas Zeitzoff. 2017. "A Bottom-Up Theory of Public Opinion about Foreign Policy." *American Journal of Political Science* 61 (3): 543–558. <https://doi.org/10.1111/ajps.12314>.
- Krosnick, Jon, and Shibley Telhami. 1995. "Public Attitudes Toward Israel: A Study of the Attentive and Issue Publics." *International Studies Quarterly* 39 (4): 535–554. <https://doi.org/10.2307/2600805>.
- Krosnick, Jon A. 1989. "Attitude Importance and Attitude Accessibility." *Personality and Social Psychology Bulletin* 15 (3): 297–308. <https://doi.org/10.1177/0146167289153002>.
- . 1990. "Government Policy and Citizen Passion: A Study of Issue Publics in Contemporary America." *Political Behavior* 12 (1): 59–92. <https://doi.org/10.1007/BF00992332>.
- . 1991. "Response strategies for coping with the cognitive demands of attitude measures in surveys." *Applied Cognitive Psychology* 5 (3): 213–236. <https://doi.org/10.1002/acp.2350050305>.

- Kupatadze, A., and T. Zeitzoff. 2021. “In the Shadow of Conflict: How Emotions, Threat Perceptions and Victimization Influence Foreign Policy Attitudes.” *British Journal of Political Science* 51 (1): 181–202. <https://doi.org/10.1017/S0007123418000479>.
- Lakens, Daniël. 2017. “Equivalence Tests: A Practical Primer for t Tests, Correlations, and Meta-Analyses.” *Social Psychological and Personality Science* 8 (4): 355–362. <https://doi.org/10.1177/194855061769717>.
- Lerner, Jennifer S., and Dacher Keltner. 2001. “Fear, anger, and risk.” *Journal of Personality and Social Psychology* (US) 81 (1): 146–159. <https://doi.org/10.1037/0022-3514.81.1.146>.
- Levin, Dov H, and Paul Musgrave. 2025. “To What End? Policy Objectives and US Public Support for Political Warfare.” *Foreign Policy Analysis* 21 (2): oraf008. <https://doi.org/10.1093/fpa/oraf008>.
- Li, Shoushi, and Luofu Ye. 2015. “How Do Americans Evaluate China’s International Responsibility? An Empirical Assessment.” *Journal of Contemporary China* 24 (92): 222–239. <https://doi.org/10.1080/10670564.2014.932149>.
- Li, Xiaojun. 2021. “More than Meets the Eye: Understanding Perceptions of China Beyond the Favorable–Unfavorable Dichotomy.” *Studies in Comparative International Development* 56 (1): 68–86. <https://doi.org/10.1007/s12116-021-09320-1>.
- Marcus, George E., W. Russell Neuman, and Michael MacKuen. 2000. *Affective Intelligence and Political Judgment*. Chicago, IL: University of Chicago Press, October.
- Martínez-Fernández, Andrés, and Andrew Harding. 2024. *Holding China and Mexico Accountable for America’s Fentanyl Crisis*, September. <https://www.heritage.org/china/report/holding-china-and-mexico-accountable-americas-fentanyl-crisis>.

- McLean, Elena V., and Dwight A. Roblyer. 2017. "Public Support for Economic Sanctions: An Experimental Analysis." *Foreign Policy Analysis* 13 (1): 233–254. <https://doi.org/10.1093/fpa/orw014>.
- Page, Benjamin I., and Robert Y. Shapiro. 1992. *The Rational Public: Fifty Years of Trends in Americans' Policy Preferences*. Chicago, IL: University of Chicago Press.
- Rathbun, Brian C., Joshua D. Kertzer, Jason Reifler, Paul Goren, and Thomas J. Scotto. 2016. "Taking Foreign Policy Personally: Personal Values and Foreign Policy Attitudes." *International Studies Quarterly* 60 (1): 124–137. Accessed March 18, 2026. <https://doi.org/10.1093/isq/sqv012>.
- Renshon, Jonathan, Julia J. Lee, and Dustin Tingley. 2017. "Emotions and the Micro-Foundations of Commitment Problems" [in en]. *International Organization* 71 (S1): S189–S218. <https://doi.org/10.1017/S0020818316000473>.
- Rossiter, Erin Leigh, and Jeffrey Harden. 2026. "The Development of an Issue Public: Evidence from The Eras Tour." *The Journal of Politics* 0 (ja): null. <https://doi.org/10.1086/737279>.
- Saunders, Heather, Nirmita Panchal, and Robin Rudowitz. 2026. *Opioid Overdose Deaths: National Trends and Variation by Demographics and States*, February. <https://www.kff.org/mental-health/opioid-overdose-deaths-national-trends-and-variation-by-demographics-and-states/>.
- Schuman, Howard, and Stanley Presser. 1981. *Questions and Answers in Attitude Surveys: Experiments on Question Form, Wording, and Context*. New York: Academic Press.
- The White House. 2025. *Imposing Duties to Address the Synthetic Opioid Supply Chain in the People's Republic of China*, February. <https://www.whitehouse.gov/presidential-actions/2025/02/imposing-duties-to-address-the-synthetic-opioid-supply-chain-in-the-peoples-republic-of-china/>.

- Tomz, Michael R., and Jessica L. P. Weeks. 2013. "Public Opinion and the Democratic Peace." *American Political Science Review* 107 (4): 849–865. <https://doi.org/10.1017/S0003055413000488>.
- U.S. House Select Committee on the Strategic Competition Between the United States and the Chinese Communist Party. 2024. *The CCP's Role in the Fentanyl Crisis*. Technical report. Washington, D.C. <https://chinaselectcommittee.house.gov/sites/evo-subsites/selectcommittteeontheccp.house.gov/files/evo-media-document/The%20CCP%27s%20Role%20in%20the%20Fentanyl%20Crisis%204.16.24%20%281%29.pdf>.
- USAFacts. 2025. *Are fentanyl overdose deaths rising in the US?*, October. <https://usafacts.org/articles/are-fentanyl-overdose-deaths-rising-in-the-us/>.
- Vangelov, Kasey, Keith Humphreys, Jonathan P. Caulkins, Harold Pollack, Bryce Pardo, and Peter Reuter. 2026. "Did the illicit fentanyl trade experience a supply shock?" *Science* 391 (6781): 134–136. <https://doi.org/10.1126/science.aea6130>.
- Wang, Donghui, Yan Ming, Shawn Dorius, and Yu Xie. 2022. "In the Eyes of the Beholder: How China and the US See Each Other." *Journal of Contemporary China* 31 (134): 232–249. <https://doi.org/10.1080/10670564.2021.1945737>.
- Wayne, Carly N. 2023. "Terrified or Enraged? Emotional Microfoundations of Public Counterterror Attitudes." *International Organization* 77 (4): 824–847. <https://doi.org/10.1017/S0020818323000152>.
- Wick, Shelley. 2014. "Capabilities, Cooperation, and Culture: Mapping American Ambivalence Toward China1." *Foreign Policy Analysis* 10 (3): 289–309. <https://doi.org/10.1111/fpa.12020>. <https://doi.org/10.1111/fpa.12020>.
- Zaller, John R. 1992. *The Nature and Origins of Mass Opinion*. Cambridge, UK: Cambridge University Press.

A Appendix: Supplementary Tables and Figures

A.1 Robustness: Alternative Blame Measure and Ordinal Models

Table 2 (in the main text) replicates the posture models using `blame_china_alt`, the single-choice “most responsible” item. The $\text{warmth} \times \text{blame}$ interaction is attenuated but retains the expected sign.

Table A1 reports ordinal logistic regression estimates treating individual posture items as ordered outcomes.

Table A1: Ordinal Logistic Regression: Individual Posture Items as Ordered Outcomes

Item	Wave	Warmth (b/SE/p)	Blame (b/SE/p)	Interaction (b/SE/p)
Defensive (a)	2024	0.345 (0.073) 0.000***	-0.172 (0.085) 0.042*	0.324 (0.097) 0.001***
Defensive (a)	2025	0.349 (0.095) 0.000***	-0.247 (0.115) 0.032*	0.462 (0.126) 0.000***
Defensive (a)	Pooled	0.342 (0.057) 0.000***	-0.197 (0.068) 0.004**	0.379 (0.076) 0.000***
Cooperative (b)	2024	0.703 (0.076) 0.000***	-0.230 (0.089) 0.010**	0.078 (0.110) 0.480
Cooperative (b)	2025	0.638 (0.095) 0.000***	-0.322 (0.109) 0.003**	0.260 (0.122) 0.033*
Cooperative (b)	Pooled	0.672 (0.059) 0.000***	-0.256 (0.067) 0.000***	0.151 (0.081) 0.063
Friendly (c, rev.)	2024	0.749 (0.073) 0.000***	0.082 (0.088) 0.353	0.210 (0.099) 0.034*
Friendly (c, rev.)	2025	0.500 (0.097) 0.000***	-0.009 (0.104) 0.935	0.284 (0.123) 0.020*
Friendly (c, rev.)	Pooled	0.639 (0.058) 0.000***	0.044 (0.067) 0.513	0.245 (0.077) 0.001**
Constructive (d)	2024	0.557 (0.070) 0.000***	0.088 (0.089) 0.319	0.096 (0.103) 0.347
Constructive (d)	2025	0.493 (0.092) 0.000***	0.142 (0.107) 0.184	0.126 (0.118) 0.287
Constructive (d)	Pooled	0.521 (0.056) 0.000***	0.115 (0.066) 0.083	0.114 (0.076) 0.133

Notes. Survey-weighted ordinal logistic regressions (`svyolr`). Each row reports a separate model for one of four posture items (scale 1–5). Coefficients reflect the log-odds of a higher response. Standard errors are design-based. n is unweighted. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A2 reports replacement robustness models substituting the behavior index for blame in the interaction term.

Table A2: Robustness: Behavior Index as Replacement for Blame Attribution

	R1: blame_china	R1: blame_china_alt
(Intercept)	0.50*** (0.08)	0.50*** (0.09)
behavior_index_z	0.34*** (0.03)	0.40*** (0.02)
blame_china	-0.07* (0.03)	
factor(year)2025	0.14*** (0.03)	0.15*** (0.03)
party3Republican	-0.08 (0.05)	-0.06 (0.05)
party3Independent	0.01 (0.04)	0.02 (0.04)
ideo5_clean	-0.10*** (0.02)	-0.11*** (0.02)
educ_college	-0.00 (0.03)	0.01 (0.03)
age	-0.00*** (0.00)	-0.00*** (0.00)
female	0.05 (0.03)	0.05 (0.03)
race_ethBlack	0.02 (0.05)	0.04 (0.05)
race_ethHispanic	0.04 (0.06)	0.03 (0.06)
race_ethAsian	-0.16 (0.09)	-0.15 (0.09)
race_ethOther	0.02 (0.05)	0.04 (0.06)
newsint_attn	0.01 (0.02)	0.01 (0.02)
exposure_index	-0.00 (0.03)	0.00 (0.03)
behavior_index_z:blame_china	0.14*** (0.03)	
blame_china_alt		-0.16*** (0.05)
behavior_index_z:blame_china_alt		0.07 (0.05)
Deviance	3095.58	2935.01
Dispersion	0.73	0.73
Num. obs.	4234	4026

*** $p < 0.001$; ** $p < 0.01$; * $p < 0.05$

Notes. Survey-weighted OLS. Dependent variable: standardized foreign-policy posture index (posture_z). Columns substitute the behavior index (behavior_index_z) for the binary blame variable as the conditioning variable. Standard errors are design-based. 40* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A3: Acquiescence Robustness: S+A Blame Contrasts Under Endorsement-Count Control

variant	Dpp_SA(w=-1)	Dpp_SA(w=+1)
Baseline	0.2944301	0.3058172
+ main + miss	0.2111030	0.1910461
+ zero-imputed	0.2111030	0.1910461

Notes. Survey-weighted multinomial logit marginal effects. Each row reports the blame-driven shift in $\hat{P}(S+A)$ at the indicated warmth level, with and without endorsement-count control for acquiescent response style. The endorsement-count covariate is the count of items the respondent endorsed across the broader policy-tool checklist, *excluding* the two items (sanctions, action) that define the S+A tool configuration to avoid mechanical correlation with the dependent variable; counts are computed within wave from the items shared across both 2024 and 2025 instruments. Higher values indicate a stronger general tendency to endorse listed items independent of any specific position on China policy. The endorsement-count control attenuates the blame effect on $P(S+A)$ by approximately 38%.




Figure A1: Placebo test: Interaction plot for sanctions $\times$ international cooperation composite as dependent variable. The placebo pairing shows a comparable blame-driven shift ($\Delta_{pp} = 0.311$ at warmth = +1 SD), indicating that the S+A effect is not fully specific to fentanyl cooperation.

Table A4: Full MNL Coefficient Estimates: Pooled, 2024, and 2025 Waves

	Pooled	2024	2025
Sanctions-only: (Intercept)	−1.89*** (0.24)	−1.77*** (0.29)	−2.35*** (0.41)
Sanctions-only: warmth_z	0.05 (0.06)	0.04 (0.08)	0.07 (0.10)
Sanctions-only: blame_china	0.80*** (0.09)	0.78*** (0.12)	0.88*** (0.15)
Sanctions-only: party3Republican	0.11 (0.13)	0.06 (0.16)	0.07 (0.23)
Sanctions-only: party3Independent	−0.27* (0.11)	−0.29* (0.14)	−0.30 (0.20)
Sanctions-only: ideo5_clean	0.16*** (0.05)	0.13* (0.06)	0.22* (0.08)
Sanctions-only: educ_college	0.21* (0.09)	0.21 (0.11)	0.22 (0.15)
Sanctions-only: age	−0.00 (0.00)	0.00 (0.00)	−0.01* (0.00)
Sanctions-only: female	−0.10 (0.08)	−0.20 (0.11)	0.09 (0.14)
Sanctions-only: race_ethBlack	0.11 (0.14)	−0.01 (0.17)	0.30 (0.22)
Sanctions-only: race_ethHispanic	0.12 (0.14)	0.17 (0.18)	0.04 (0.23)
Sanctions-only: race_ethAsian	0.41 (0.23)	0.12 (0.29)	0.93* (0.38)
Sanctions-only: race_ethOther	−0.27 (0.19)	−0.08 (0.23)	−0.67* (0.33)
Sanctions-only: newsint_attn	0.13** (0.05)	0.10 (0.06)	0.17* (0.08)
Sanctions-only: exposure_index	0.13 (0.08)	0.06 (0.10)	0.23* (0.12)
Sanctions-only: posture_z	−0.13** (0.05)	−0.12* (0.06)	−0.15 (0.08)
Sanctions-only: factor(year)2025	−0.24** (0.09)		
Sanctions-only: warmth_z:blame_china	−0.19* (0.09)	−0.15 (0.12)	−0.25 (0.14)
Action-only: (Intercept)	−2.52*** (0.33)	−2.28*** (0.41)	−3.10*** (0.56)
Action-only: warmth_z	0.06 (0.10)	0.17 (0.13)	−0.05 (0.17)
Action-only: blame_china	1.65*** (0.13)	1.56*** (0.17)	1.84*** (0.21)
Action-only: party3Republican	0.22 (0.19)	0.18 (0.23)	0.32 (0.32)
Action-only: party3Independent	0.04 (0.16)	−0.24 (0.21)	0.38 (0.25)
Action-only: ideo5_clean	−0.07 (0.07)	−0.04 (0.08)	−0.14 (0.11)
Action-only: educ_college	0.01 (0.13)	−0.28 (0.17)	0.37 (0.19)

Table A5: Design-Based Sensitivity Check: MNL S+A Coefficients (Log-Odds Scale)

Outcome	Predictor	Estimate	Design SE	t-value	p-value	Stars
Sanctions-only	warmth_z	0.0474	0.0619	0.766	0.4439	
Action-only	warmth_z	0.0632	0.1016	0.622	0.5342	
S+A	warmth_z	-0.2269	0.1018	-2.228	0.0259	*
Sanctions-only	blame_china	0.8045	0.0927	8.676	0.0000	***
Action-only	blame_china	1.6519	0.1314	12.572	0.0000	***
S+A	blame_china	2.6398	0.1130	23.353	0.0000	***
Sanctions-only	factor(year)2025	-0.2426	0.0871	-2.785	0.0054	**
Action-only	factor(year)2025	-0.0559	0.1191	-0.469	0.6391	
S+A	factor(year)2025	-0.1267	0.0922	-1.373	0.1697	
Sanctions-only	warmth_z:blame_china	-0.1927	0.0894	-2.154	0.0313	*
Action-only	warmth_z:blame_china	-0.1546	0.1257	-1.230	0.2189	
S+A	warmth_z:blame_china	0.1343	0.1143	1.175	0.2402	

Notes. Design-based standard errors computed via `svyglm()` binary contrasts (S+A vs. all other categories) on the log-odds scale. The binary-contrast approach approximates the design-based equivalent of the full MNL. Point estimates are consistent with the main multinomial specification. The blame main effect on S+A is robust ($b = 2.640$, $p < 0.001$); the $\text{warmth} \times \text{blame}$ interaction is not statistically distinguishable from zero under design-based inference ($b = 0.134$, $p = 0.240$).

Table A6: Design-Based vs. Model-Based Standard Errors on the Probability Scale: Blame Contrast $\Delta_{pp}(\text{S+A})$

Warmth	MNL Δ_{pp}	MNL Hessian SE	svyglm Design SE	SE Ratio
-1 SD	0.3049	1.2730	0.0188	0.015
+0 SD	0.3160	1.2780	0.0140	0.011
+1 SD	0.3239	1.2749	0.0229	0.018

Notes. Design-based SEs computed from `svyglm()` binary contrasts (S+A vs. other) and converted to the probability scale via the standard delta method. MNL Hessian SEs are computed from the full multinomial model's variance-covariance matrix (also delta-method propagated). The Hessian SEs are an order of magnitude larger than the design-based SEs at every warmth level, reflecting that the full-MNL Hessian variance for predicted probabilities incorporates correlation across alternatives that magnifies uncertainty for $\hat{P}(\text{S+A})$. The design-based SEs from the binary-contrast specification are the recommended quantities for inference about $\Delta_{pp}(\text{S+A})$. Point estimates of Δ_{pp} differ slightly from those in Table 3 because the binary contrast is fit on $\mathbf{1}\{\text{tool} = \text{S+A}\}$ rather than on the full four-category outcome. SE Ratio = `svyglm Design SE` / MNL Hessian SE.

Table A7: Eight-Probe Panorama: S+A Minus Sanctions-Only Contrasts Across All Scenarios

Probe	Domain	Estimate	SE	CI	p_value
a: Taiwan military action	Security	-0.0750	0.0688	[-0.2099, 0.0599]	0.276
b: US-China military collision	Security	0.0285	0.0826	[-0.1335, 0.1905]	0.730
c: Other military collision	Security	0.0728	0.0734	[-0.0711, 0.2168]	0.321
d: Fentanyl cooperation	Health	0.2817	0.0913	[0.1027, 0.4607]	0.002**
e: Resuming adoptions	Cultural	0.0837	0.0872	[-0.0873, 0.2547]	0.337
f: Electronics sales stopped	Economic	0.1257	0.0916	[-0.0540, 0.3053]	0.170
g: Student scholarships	Cultural	0.0160	0.0882	[-0.1570, 0.1891]	0.856
h: Pandas to US zoos	Cultural	0.0402	0.0769	[-0.1106, 0.1911]	0.601

Notes. Survey-weighted stacked OLS. Each row reports the blame-driven S+A minus Sanctions-only contrast (Δ_s) for one of eight probes, among warm respondents ($warmth_z = +1$ SD). Standard errors are clustered at the respondent level. 95% confidence intervals reported. p -values are two-sided.

Table A8: Pairwise Divergence Tests: Target Probe vs. Boundary Probes (BH-FDR Corrected)

Comparison	Divergence	SE	CI	p_value	p_fdr
d vs a: Taiwan military action	0.3567	0.1200	[0.1213, 0.5921]	0.003**	0.021*
d vs b: US-China military collision	0.2532	0.1208	[0.0163, 0.4902]	0.036*	0.063
d vs c: Other military collision	0.2089	0.1122	[-0.0111, 0.4289]	0.063	0.088
d vs e: Resuming adoptions	0.1980	0.1117	[-0.0210, 0.4170]	0.076	0.089
d vs f: Electronics sales stopped	0.1560	0.1202	[-0.0798, 0.3919]	0.195	0.195
d vs g: Student scholarships	0.2657	0.1048	[0.0601, 0.4712]	0.011*	0.040*
d vs h: Pandas to US zoos	0.2415	0.1089	[0.0279, 0.4550]	0.027*	0.062

Notes. Each row reports a pairwise contrast between Δ_d (probe d , fentanyl cooperation) and one boundary probe. p -values are two-sided; p_{FDR} applies the Benjamini–Hochberg false discovery rate correction across seven comparisons. Standard errors are clustered at the respondent level.

Table A9: Alternative Specification: Zero-Imputed Endorsement Counts

Probe	Estimate	SE	CI_low	CI_high	p
probe_d	0.2816971	0.0912776	0.1026776	0.4607166	0.0020581

Notes. Blame-driven S+A minus Sanctions-only contrasts using zero-imputed endorsement counts (non-responses coded 0 rather than excluded). Compare to main estimates in Table 5.

Table A10: Alternative Specification: Excluding Non-Respondents

Probe	Estimate	SE	CI_low	CI_high	p
probe_d	0.2388492	0.0994101	0.0438798	0.4338186	0.0163753

Notes. Blame-driven S+A minus Sanctions-only contrast for probe d (fentanyl cooperation) among respondents who answered all relevant items (complete-case exclusion). Compare to main estimate in Table 5.

Table A11: Alternative Specification: Continuous $\hat{P}(S+A)$ as Dependent Variable


DV	Estimate	SE	t_value	p_value	Note
probe_d	0.5018874	0.9951674	0.5043246	0.6140597	SEs do not propagate Stage 1 uncertainty; directional only

Notes. Blame-driven shifts in MNL-predicted $\hat{P}(S+A)$ (continuous, replacing the binary S+A classifier). Compare to main estimates in Table 5.

Table A12: TOST Equivalence Test: Sensitivity to Equivalence Bound Choice

Bound ($\times$ SD)	ε	Composite $\bar{\Delta}$	SE	TOST p	Reject H_0
0.10	0.145	0.042	0.043	0.009	Yes
0.15	0.218	0.042	0.043	<0.001	Yes
0.20	0.290	0.042	0.043	<0.001	Yes
0.25	0.363	0.042	0.043	<0.001	Yes

Notes. Each row reports the TOST result at a different equivalence bound (expressed as a fraction of the boundary-probe SD). The composite boundary estimate ($\bar{\Delta} = 0.042$, $SE = 0.043$) is negligibly small at all four bound choices. H_0 : the composite effect exceeds the bound in absolute value; rejection of H_0 implies the effect is equivalently small.

**Figure A2:** Coefficient forest plot for the warmth $\times$ blame interaction across model specifications. Point estimates with 95% confidence intervals.

A.2 Robustness: Acquiescence and Full Coefficients

A.3 Robustness: Full Panorama, Pairwise Divergence, and Alternative Specifications

A.4 Cross-Wave Robustness: Interaction Stability Across Survey Years

The pooled models in the main text include year fixed effects to account for mean-level differences across the 2024 and 2025 waves. Table A13 tests whether the warmth $\times$ blame interaction itself changed between waves by estimating a triple interaction term (warmth $\times$ blame $\times$ year). The non-significant triple term ($b = 0.068$, $p = 0.332$) provides no evidence that the interaction shifted between the two waves. Table A14 reports wave-specific estimates confirming that the warmth $\times$ blame interaction is statistically significant in both the 2024 and 2025 waves individually.

Table A13: Cross-Wave Stability: Warmth $\times$ Blame $\times$ Year Triple Interaction

Quantity	Estimate	SE	95% CI	p-value	N
2024 interaction (warmth x blame)	0.1039	0.0461	[0.0134, 0.1943]	0.024*	4234
2025 - 2024 differential (triple term)	0.0679	0.0700	[-0.0692, 0.2051]	0.332	4234
2025 interaction (implied sum)	0.1718	0.0531	[0.0678, 0.2758]	0.001**	4234

Notes. Survey-weighted OLS. The 2025–2024 differential (triple term) tests whether the warmth $\times$ blame interaction changes across waves. The non-significant triple term ($b = 0.068$, $p = 0.332$) supports cross-wave stability. Implied 2025 interaction is the sum of the 2024 coefficient and the differential.

Table A14: Within-Wave Replication: Warmth $\times$ Blame Interaction by Survey Year

Wave	N	Warmth (b/SE)	Blame (b/SE)	Interaction (b)	Interaction (SE)	Interaction (p)	Inter
2024	2545	0.3862 (0.0331)	-0.0526 (0.0418)	0.0961	0.0461	0.037*	[0.0
2025	1689	0.3032 (0.0395)	-0.0694 (0.0500)	0.1846	0.0527	0.000***	[0.0
Pooled (year FE)	4234	0.3486 (0.0255)	-0.0532 (0.0319)	0.1340	0.0347	0.000***	[0.0

Notes. Survey-weighted OLS estimated separately for each wave and pooled with year fixed effects. The warmth $\times$ blame interaction is statistically significant in both waves individually ($b_{2024} = 0.096$, $p = 0.037$; $b_{2025} = 0.185$, $p < 0.001$). The 2025 coefficient is larger, consistent with the post-tariff political context heightening the conditional pattern.

A.5 Placebo: Sanctions and International Cooperation

A.6 Post-Treatment Sensitivity: Total-Effect MNL Specification

The main multinomial logit specification controls for `posture_z`, which is a post-treatment covariate with respect to warmth and blame. Conditioning on a post-treatment variable can, in principle, bias the estimated effect of the primary predictors by absorbing part of the treatment path. Table A16 compares the within-posture specification (`posture_z` included) against a total-effect specification that omits `posture_z` entirely. The blame-driven shift in $\hat{P}(S+A)$ is $\Delta = 0.306$ ($SE = 0.029$) under the within-posture specification and $\Delta = 0.305$ ($SE = 0.028$) under the total-effect specification at $\text{warmth} = +1$ SD; the difference between the two contrasts is less than 0.001. Post-treatment conditioning does not materially alter the estimated blame effect on tool configuration, confirming that the within-posture estimates carry the same substantive interpretation as the total-effect estimates. Standard errors in this table are Hessian-based (model-based) and are not directly comparable to design-based standard errors reported in the main results.

A.7 Formal Placebo Contrast: Main vs. Placebo Blame Effect

Section 5 notes that the placebo MNL (using $\text{sanctions} \times \text{international cooperation}$ as the dependent variable) produces a blame-driven shift comparable in magnitude to the main specification. Table A17 formalizes this comparison. The main blame contrast is $\Delta_d = 0.282$ ($SE = 0.091$), derived from the eight-probe specificity model (probe d , fentanyl cooperation, among $S+A$ and Sanctions-only respondents). The placebo blame contrast— $\Delta_{pp}(S+A)$ at $\text{warmth} = +1$ SD from the $\text{sanctions} \times \text{international cooperation}$ model—is 0.311 ($SE = 0.027$). The difference between the two contrasts is -0.029 ($SE = 0.095$, $t = -0.31$, $p = 0.758$, 95% CI $[-0.216, 0.157]$). The two estimates are statistically indistinguishable, consistent with both configurations being engagement-compatible pressure packages that are similarly activated by blame attribution. This result supports the domain-concentration interpretation of issue-bounded conditionality (IBC) but does not support a stronger tool-specificity interpretation. The combined standard error in this table assumes inde-

Table A15: Placebo MNL: Sanctions $\times$ International Cooperation as Dependent Variable

	Placebo (Pooled)
Sanctions-only: (Intercept)	−2.11*** (0.26)
Sanctions-only: warmth_z	0.07 (0.07)
Sanctions-only: blame_china	0.96*** (0.10)
Sanctions-only: party3Republican	0.21 (0.14)
Sanctions-only: party3Independent	−0.08 (0.13)
Sanctions-only: ideo5_clean	0.13* (0.05)
Sanctions-only: educ_college	0.21* (0.10)
Sanctions-only: age	0.00 (0.00)
Sanctions-only: female	−0.18* (0.09)
Sanctions-only: race_ethBlack	0.15 (0.15)
Sanctions-only: race_ethHispanic	0.12 (0.15)
Sanctions-only: race_ethAsian	0.38 (0.26)
Sanctions-only: race_ethOther	−0.18 (0.21)
Sanctions-only: newsint_attn	0.10 (0.05)
Sanctions-only: exposure_index	0.10 (0.08)
Sanctions-only: posture_z	−0.10 (0.05)
Sanctions-only: factor(year)2025	−0.17 (0.09)
Sanctions-only: warmth_z:blame_china	−0.21* (0.10)
IntlCoop-only: (Intercept)	−1.35*** (0.26)
IntlCoop-only: warmth_z	−0.02 (0.07)
IntlCoop-only: blame_china	0.68*** (0.10)
IntlCoop-only: party3Republican	0.16 (0.15)
IntlCoop-only: party3Independent	0.26* (0.12)
IntlCoop-only: ideo5_clean	−0.32*** (0.05)
IntlCoop-only: educ_college	0.17 (0.10)

Table A16: Post-Treatment Sensitivity: Within-Posture vs. Total-Effect MNL Specification

Specification	Outcome	Warmth x Blame	SE (model-based)
— MNL Coefficients (Warmth x Blame interaction) —			
Within-posture (posture_z included)	Sanctions-only	-0.1927	0.0894
Total-effect (posture_z excluded)	Sanctions-only	-0.2098	0.0893
Within-posture (posture_z included)	Action-only	-0.1546	0.1257
Total-effect (posture_z excluded)	Action-only	-0.1578	0.1254
Within-posture (posture_z included)	S+A	0.1343	0.1143
Total-effect (posture_z excluded)	S+A	0.1289	0.1142
— Blame contrasts: Dpp(S+A) at warmth levels —			
Within-posture (posture_z included)	warmth_z = -1	0.2944 (SE=0.0248)	(predicted)
Total-effect (posture_z excluded)	warmth_z = -1	0.2938 (SE=0.0246)	(predicted)
Within-posture (posture_z included)	warmth_z = +0	0.3019 (SE=0.0234)	(predicted)
Total-effect (posture_z excluded)	warmth_z = +0	0.3018 (SE=0.0233)	(predicted)
Within-posture (posture_z included)	warmth_z = +1	0.3058 (SE=0.0289)	(predicted)
Total-effect (posture_z excluded)	warmth_z = +1	0.3051 (SE=0.0284)	(predicted)
Note: MNL SEs are Hessian-based (model-based); not directly comparable to design-based SEs in main tables.			

Notes. Each panel compares the within-posture specification (posture_z included as a control) against the total-effect specification (posture_z omitted). The upper panel reports MNL log-odds coefficients on the warmth × blame interaction; the lower panel reports predicted probability contrasts (Δ_{pp}) for the S+A outcome at three warmth levels. Standard errors are Hessian-based (model-based) and are not directly comparable to design-based standard errors in the main tables. The reference category is Neither.

pendence of the two models; because respondents overlap across both estimation samples, this assumption is anti-conservative; the reported p -value understates uncertainty and should be treated as a lower bound.

A.8 Independence of Irrelevant Alternatives: Formal Tests and Heuristic

The multinomial logit assumes that the relative odds between any two alternatives are independent of the presence of other alternatives (IIA). Table A18 reports Hausman-McFadden formal tests for each alternative. IIA is rejected for the Neither and S+A categories ($\chi^2 = 1017.6$ and $\chi^2 = 1257.5$ respectively, both $p < 0.001$). It is not rejected for Sanctions-only ($\chi^2 = 19.4$, $p = 0.989$) or Action-only ($\chi^2 = -3.5$, $p = 1.000$). The negative chi-square value for Action-only arises when the restricted-model variance-covariance matrix is not positive definite, a known property of the Hausman-McFadden procedure that indicates no evidence against IIA for that alternative.

The substantive content of an IIA rejection is often misread, so I state the implication explicitly. Rejection means that removing one alternative from the choice set would redistribute its probability mass *non-proportionally* across the remaining alternatives, rather than maintaining the original

Table A17: Formal Placebo Contrast: Main vs. Placebo Blame Effects on S+A Probability

Estimand	Value
Δ_d (probe d : fentanyl cooperation)	0.2817 ($SE = 0.0913$)
$\Delta_{pp,placebo}$ (Sanctions+IntlCoop, $w = +1$)	0.3110 ($SE = 0.0268$)
Difference: $\Delta_d - \Delta_{pp,placebo}$	-0.0293
Standard error (independence assumption)	0.0951
t -statistic	-0.3079
p -value (two-sided)	0.758
95% CI lower	-0.2158
95% CI upper	0.1573
Approximate df (Welch-Satterthwaite)	1973.1

Notes. Δ_d is the S+A-minus-S-only adjusted mean contrast for probe d (fentanyl cooperation) from the eight-probe specificity model (Table 5). $\Delta_{pp,placebo}$ is the predicted-probability blame contrast for the sanctions $\times$ international cooperation outcome at warmth = +1 SD from the placebo multinomial logit. The combined SE assumes independence of the two models; overlap in respondent samples may understate true uncertainty. Degrees of freedom use the Welch-Satterthwaite approximation.

relative odds. It does not mean the multinomial logit estimates are invalid, nor does it imply that the categories are “distinctive” as preference objects in any deeper sense. I therefore interpret the MNL coefficients and predicted probabilities reported in the main text as conditional associations *within the four-alternative menu actually presented to respondents*—i.e., as descriptions of how blame and warmth co-vary with the chosen tool configuration in this particular choice environment—rather than as structural preferences that would extrapolate to a different option set. The coefficient-stability heuristic in Table A19 provides a complementary check: estimates of the warmth, blame, and warmth-by-blame coefficients shift by less than 0.10 when the S+A category is omitted, supporting the use of MNL results for within-menu inference. The rejection pattern is also consistent with the theoretical framing of Section 2.2: S+A and Neither occupy distinct attitudinal regions in the issue-bounded conditionality logic, whereas Sanctions-only and Action-only behave as closer substitutes within the punitive-cooperative space. I do not estimate a nested-logit alternative because the unweighted-only mlogit fit failed to converge in this sample and because the conceptual clarification, rather than an alternative discrete-choice estimator, is the appropriate response to an IIA rejection in this design.

As a supplementary heuristic, Table A19 reports the maximum coefficient shift observed when

Table A18: IIA Formal Tests: Hausman-McFadden Statistic for Each Alternative

Dropped alternative	χ^2	df	p -value	Verdict
Neither	1017.559	36	0.0000	Reject IIA (alt has distinctive effects)
Sanctions-only	19.402	36	0.9892	Fail to reject IIA
Action-only	-3.461	36	1.0000	Fail to reject IIA
S+A	1257.505	36	0.0000	Reject IIA (alt has distinctive effects)

Notes. Hausman-McFadden IIA test. Each row tests IIA by dropping one alternative from the choice set and comparing restricted versus full-model coefficient estimates. χ^2 values with negative sign indicate that the variance-covariance difference matrix is not positive definite, interpreted as no evidence against IIA. p -values are computed from the χ^2 distribution with $df = 36$.

the S+A category is omitted from the choice set and the remaining three-category model is re-estimated. The maximum shift across all coefficients is 0.034, well below the conventional 0.10 stability threshold. This result is consistent with the formal test conclusions.

Table A19: IIA Heuristic: Coefficient Stability After Omitting S+A Category

Item	Result
Method	Coefficient stability comparison (mlogit unavailable)
Max coef. shift (heuristic)	0.0335
Degrees of freedom	—
p -value	—
Verdict	Stable (max shift ≤ 0.10); formal IIA test not run
Note	Heuristic proxy only; formal Hausman-McFadden test requires mlogit package.

Notes. The heuristic compares MNL coefficient estimates from the full four-category model against estimates from a restricted model that omits the S+A category. The maximum absolute coefficient shift is the primary diagnostic. A shift below 0.10 is interpreted as consistent with IIA. See Table A18 for the formal Hausman-McFadden tests.

A.9 Multiple Imputation Robustness Check

Listwise deletion on ideology (`ideo5_clean`), the primary source of missing data, reduces the analytic sample from 5,000 to 4,234 respondents. To assess whether this selective attrition materially affects the main interaction estimate, Table A20 compares the complete-case estimate against a multiple imputation (MI) estimate obtained using Predictive Mean Matching with $m = 10$ imputations (`mice` package). The MI procedure uses frequency-weighted `lm()` as an approximation

to `svyglm()` because `svyglm` is not directly compatible with the `mice` pooling rules; the design-based standard error from the complete-case `svyglm` model remains the authoritative specification. The complete-case estimate is $b = 0.134$ ($SE = 0.035$); the MI-pooled estimate is $b = 0.123$ ($SE = 0.026$), an attenuation of approximately 8%. The fraction of missing information (FMI) is 0.001, indicating that missing data contributes negligible additional uncertainty once imputed. The missing-at-random (MAR) assumption required for valid MI is plausible given that missingness is concentrated in a single covariate and the imputation model includes all observed predictors. These results show that the complete-case $warmth \times blame$ interaction is not an artifact of listwise deletion.

Appendix Table A21 reports pooled MI estimates for the multinomial logit predicted probabilities; Appendix Table A22 reports the eight-probe specificity contrasts under full-sample imputation; Appendix Table A23 reports Rubin’s rules variance components confirming negligible FMI across all key estimates.

Table A20: Multiple Imputation Robustness: Warmth $\times$ Blame Interaction Coefficient

Specification	Warmth x Blame	SE	p-value	N	Note
Complete-case (apool, <code>svyglm</code>)	0.1340	0.0347	0.000***	4234	Design-based SE (<code>svyglm</code>); authoritative
Multiple imputation (m=10, weighted <code>lm</code>)	0.1232	0.0258	0.000***	5000	FMI=0.001; frequency-weighted <code>lm</code> ap

Notes. The complete-case estimate uses survey-weighted `svyglm()` with design-based standard errors; this is the authoritative specification. The MI estimate uses frequency-weighted `lm()` pooled across $m = 10$ imputations via Predictive Mean Matching (`mice`). FMI = fraction of missing information. $N = 4,234$ (complete-case); $N = 5,000$ (MI, full analytic sample).

Table A21: Multiple Imputation Robustness: Multinomial Logit Predicted Probabilities

Parameter	Complete case				MI ($m = 10$)			
	Est.	SE	p	N	Est.	SE	p	FMI
S+A: warmth $\times$ blame (log-odds)	0.129	0.114	0.259	4234	0.161	0.109	0.141	0.000
S+A: blame effect at warmth= +1	2.771	0.164	<0.001	4234	2.853	0.154	<0.001	0.000
S+A: warmth main effect	-0.247	0.100	0.014	4234	-0.278	0.097	0.004	0.000
S-only: warmth $\times$ blame (log-odds)	-0.210	0.089	0.019	4234	-0.171	0.083	0.039	0.001
S-only: blame effect at warmth= +1	0.600	0.134	<0.001	4234	0.682	0.122	<0.001	0.001

Notes. Pooled MI estimates ($m = 10$ imputations, PMM). Predicted probabilities for the S+A outcome at warmth = ± 1 SD under complete-case and MI specifications. FMI = fraction of missing information.

Table A22: Multiple Imputation Robustness: Eight-Probe Specificity Contrasts

Probe	Scenario	Complete case			MI pooled (Approach B)		
		Est.	SE	p	Est.	SE	FMI
a	Taiwan military action	-0.0750	0.0688	0.276	-0.1980	0.0640	0.002
b	US-China military collision	0.0285	0.0826	0.730	-0.0888	0.0724	0.002
c	Other military collision	0.0728	0.0734	0.321	-0.0703	0.0674	0.003
d†	Fentanyl cooperation (TARGET)	0.2817	0.0913	0.002	0.3143	0.0777	0.001
e	Resuming adoptions	0.0837	0.0872	0.337	-0.0535	0.0766	0.001
f	Electronics sales stopped	0.1257	0.0916	0.170	0.1131	0.0810	0.001
g	Student scholarships	0.0160	0.0882	0.856	-0.0595	0.0793	0.002
h	Pandas to US zoos	0.0402	0.0769	0.601	-0.0516	0.0716	0.001

Notes. Pooled MI estimates ($m = 10$ imputations, PMM) for the stacked-OLS probe contrasts using Approach B (impute predictors only; recompute the binary $\mathbf{1}\{\text{tool} = \text{S+A}\}$ outcome from the multinomial-logit predicted probabilities within each imputation). This approach was adopted after the original MI implementation produced no estimates for individual probe contrasts because the imputed binary outcome was a constant (all 0 or all 1) within several imputations, causing `lm()` to drop the predictor. Approach B avoids this degeneracy and yields populated estimates for all eight probes. The target probe (*d*, fentanyl cooperation) shows the largest contrast under both complete-case and MI specifications. FMI = fraction of missing information.

Table A23: Multiple Imputation Diagnostics: Variance Components and FMI

Parameter	m	V_W	V_B	T	FMI	df (capped)	Es
Phase 3: warmth_z x blame_china (posture_z)	10	0.00066	0.00000	0.00067	0.003	4214	0.122
Phase 4: S+A warmth_z x blame_china (ln-odds)	10	0.01194	0.00000	0.01194	0.000	4154	0.160
Phase 4: blame effect at warmth=+1 on S+A ln-odds	10	0.02369	0.00000	0.02369	0.000	4154	2.853
Phase 4: S-only warmth_z x blame_china (ln-odds)	10	0.00685	0.00001	0.00686	0.001	4154	-0.171
Phase 4: S-only blame effect at warmth=+1	10	0.01475	0.00001	0.01476	0.001	4154	0.681
Phase 5: Delta_d (target probe)	10	0.00603	0.00001	0.00604	0.001	2056	0.314
Phase 5: Composite boundary contrast	10	0.00077	0.00000	0.00077	0.001	2056	-0.058
Phase 5: Omnibus divergence (Delta_d - composite)	10	0.00680	0.00001	0.00681	0.001	2056	0.371

Notes. Rubin's rules variance decomposition. V_W = within-imputation variance; V_B = between-imputation variance;

FMI = fraction of missing information. Reported degrees of freedom apply the Barnard–Rubin small-sample correction with the cap at the complete-data df ($n - k$): 4,214 for Phase 3, 4,154 for Phase 4, 2,056 for Phase 5. The cap prevents the textbook Barnard–Rubin formula from returning implausibly large df (in the millions) when between-imputation variance V_B is near zero, which occurs here because the missingness pattern is concentrated in a single covariate (ideology) and the predictive-mean-matching imputer recovers nearly identical parameter estimates across imputations. Low FMI (≈ 0.001) confirms that missingness contributes negligible uncertainty to the pooled estimates.